MAD MAZE GENERATION AND DESIGN

by

Johnathon Henke

A thesis submitted to the Faculty and the Board of Trustees of the Colorado School of Mines in partial fulfillment of the requirements for the degree of Master of Science (Computer Science).

Golden, Colorado

Date _____

Signed: _____

Johnathon Henke

Signed: _____

Dr. Dinesh Mehta
Thesis Advisor

Golden, Colorado

Date _____

Signed: _____

Dr. Iris Bahar
Professor and Head
Department of Computer Science

ABSTRACT

Several logic mazes and their state graph representations are defined with the goal of generating additional instances of each maze. As Local Search proved to be an effective mad maze generation method, we focused on defining an objective function that considers maze characteristics. By representing logic mazes as state graphs, it becomes possible to employ the same objective function for scoring, generating, and contrasting various maze types. This not only simplifies implementation, but also enables the creation of common characteristic state graphs for completely unique and separate high-level puzzles. We outline desired maze features, provide implementation details, and present example objective functions for generating mazes with the targeted attributes.

TABLE OF CONTENTS

IST OF FIGURES

IST OF TABLES

CHAPTER 1

INTRODUCTION

1.1   Mazes

Mazes have long fascinated humanity. In one of the world's first civilizations in Sumer, we have discovered mazes and labyrinths. Millennia ago, isolated civilizations resulted in unique alphabets, languages, and cultures, yet mazes and labyrinth designs existed across many civilizations: from rock carvings in Sardinia (2500-2000 BC) to Padugula in South India (circa 1000 BC) to Val Camonica in Italy (750-550 BC) (Fisher and Gerster, 2000, 12). Mazes and labyrinths have inspired stories and legends, such as Daedalus, Theseus, and the legendary labyrinth in Greek mythology. The prevalence of mazes across history speaks to their extraordinary presence: "Once seen, a maze cannot be ignored. It draws you into it like a magnet, then proceeds to puzzle, infuriate, and delight in turn until its goal is reached. Mazes have been exerting this maddening fascination for thousands of years, and evidence of them is to be found in different civilizations all over the world (Fisher and Gerster, 2000, 12)."

This work explores a newer development in the history of mazes: the *Mad* Maze, also known as a logic maze.

1.2   "Mad" Mazes

The term "Mad" maze was coined by Abbott in his 1990 book of the same name. A more widespread and common term is logic mazes (Abbott, 1999, 1). These mazes introduce additional rules to become more like logic puzzles. They are also known as "Multi-state" mazes because it is possible to revisit the same location multiple times in different "states" (in many logic mazes, it is required to do so).

Consider the first maze in Abbott's *Mad Mazes*. It is a classic maze with two additional rules: after passing through a red location, one must pass through a blue location and alternate colors until reaching the exit, and one cannot turn around in place and retrace one's steps.

Figure 1.1: A maze with a similar structure as maze #1 of *Mad Mazes* (Abbott, 1990, 5).

The lower loop (highlighted above) that alternates red and blue is the defining feature. Despite appearing superfluous, it is actually required to pass through this lower loop to reach the finish. Abbott noted this when designing the maze (Abbott, 1990, 38). Adding such rules turns classic mazes into logic puzzles and makes them much more engaging. Logic mazes such as these can be modeled as graphs and solved, which leads to the purpose and motivation of this work.

1.3 Preview

Chapter 2 provides the background and motivation for this work, as well as state graph models for some logic mazes. Chapter 3 discusses several maze characteristics that can be used in an objective function to rate state graphs and high-level instances of each logic maze. Chapter 4 introduces implementation details, search strategies such as simulated annealing and stochastic local search, and an example objective function to generate mazes with specific attributes. Chapter 5 presents puzzle instance terms added to the state graph objective function for each maze along with the resulting maze instances generated. In Chapter 6, we summarize the paper and conclude with potential future work.

CHAPTER 2

BACKGROUND

2.1 Mad Mazes Project

Graph theory is a fundamental component of the undergraduate Algorithms class at Colorado School of Mines, and one of the most critical concepts covered is "graph modeling", which involves taking an unknown problem that one does not know how to solve, modeling it as a graph, and then applying a well-known graph algorithm (ideally without modifications) to solve the problem. For instance, consider the following example problem:

You are provided a set of movies $M_1, M_2, \ldots M_k$ and a set of customers, each of whom indicates two movies they would like to see this weekend. Movies are shown on Saturday and Sunday, and multiple movies may be screened at the same time. You must decide which movies should be televised on Saturday and which on Sunday, so that every customer gets to see the two movies they desire. Is there a schedule where each movie is shown at most once? Design an efficient algorithm to find such a schedule if one exists (Skiena, 2008, 188).

This problem can be solved by modeling it as a graph. Model movies as vertices and preferences as edges. If a person wants to watch movies M_1 and M_2, let there be an edge (undirected, unweighted) between M1 and M2. After creating this graph, attempt a 2-coloring using a BFS.

Graph modeling is a crucial concept, which led to the development of the Mad Maze project. While teaching the class, Dr. Mehta came across Robert Abbott's book, *Mad Mazes*. This work (which appears to have been designed as a puzzle book) contains 20 logic mazes that serve as an excellent tool to teach graph modeling, and thus the Maze project was born. Students are given a maze from the book with its additional rules and complexity beyond a classic maze, asked to model it as a graph, and run a well-known graph algorithm completely unmodified (students are encouraged to use a graph library to

3

this end) to solve the maze, which equates to finding the shortest path from the start to the finish. However, Abbott provides only one instance of each maze in the book, and solving Abbott's instance by hand is a grueling task. From a pedagogical perspective, it would be better to have smaller instances of each maze to introduce students to the rules, and additional instances to test each student's implementation, which was the direct motivation for this work: how can one generate logic mazes of a desired size and difficulty?

2.2 Modeling

As mentioned in the previous section, all mazes (classic and logic) can be modeled as a state graph where each vertex represents a state and each edge denotes a possible transition between two states. In classic mazes, one's location is the only state that needs to be tracked, and there are several methods to track the state and create vertices. One method involves modeling each "decision point" (typically intersections) as a vertex and corridors between the decisions points as edges connecting the vertices (Fig. 2.1). Classic mazes are undirected—it is always possible to retrace one's steps. However, this is often not the case with logic mazes.

The state graph representation abstracts away the visual structure of the maze. A maze's shape can impact its difficulty and confuse its solvers, especially with life-size mazes such as those found in theme parks. However, in this work, we focus on the structural elements of the maze encoded into its state graph. The physical shape of the maze and its effect on the human psyche fall outside the scope of this work.

2.2.1 Classic Maze State Graphs

As stated previously, classic mazes can be solved by modeling them as a graph and using a BFS to find the shortest path through the maze. The shortest path through the maze illustrated in Fig. 2.1 is S, 1, 3, 5, 8, 10, 12, F. This path is a simple path because it does not repeat any vertices.
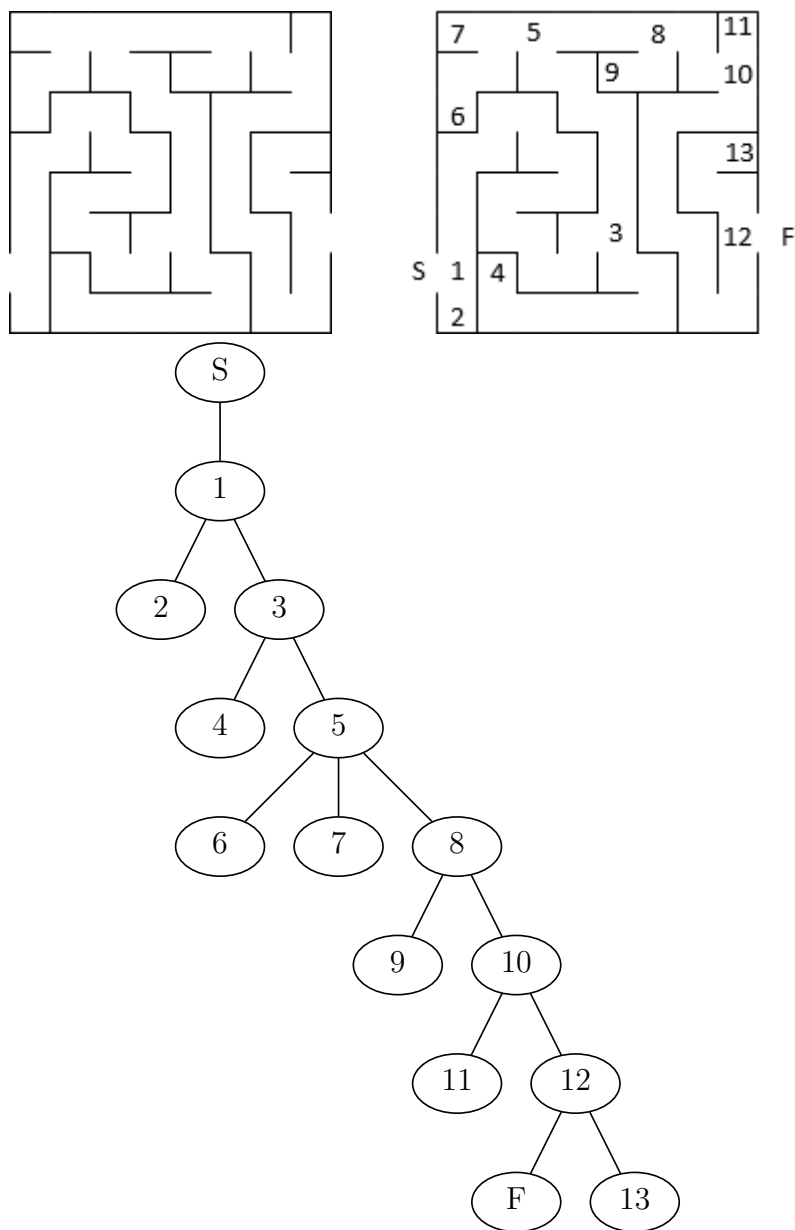
Figure 2.1: A classic maze and one possible graph representation.

2.2.2 Jumping Mazes and Variants

Jumping mazes are a common type of logic maze. The first jumping maze Abbott introduces is maze #7 Jumping Jim (Abbott, 1990, 14-15)—this simple version of a jumping maze consists of a simple numeric grid. The grid has a number on each cell that indicates how far one must move (horizontally or vertically, but not diagonally) from that cell. Starting from the upper left hand corner, the goal is to reach the bottom right hand corner in the minimum number of "jumps" (note it is the number of jumps that we try to minimize and not the total length of the jumps).

3	2	2	2
1	2	2	2
3	2	1	3
1	1	3	GOAL

Figure 2.2: A basic jumping maze.

The classic jumping maze has a simple model. Create an unweighted, directed graph G. Add vertices representing each cell on the grid and directed edges from each vertex to the cells that can be reached. For example, the vertex for (1, 1) in the grid shown in Figure 2.2 would have directed edges to (4, 1) and (1, 4). Since the graph is unweighted, a BFS is sufficient to find the shortest path from the start to the finish, and there is no need to use Dijkstra's or another weighted graph algorithm.

Due to this puzzle's simple state graph reduction, classic jumping mazes were not used for the graph modeling project. Instead, the jumping maze variant introduced in maze #15 Jumping Jim's Encore was used (Abbott, 1990, 26-27). This variant introduces circled numbers that change the direction of movement. If one lands on a circled number while moving horizontally or vertically, then one's movement direction changes to diagonal until

reaching another circled number, at which point it reverts to vertical/horizontal. This small change introduces a concept of "movement state" into the maze, making it possible to visit the same cell multiple times, once in the state of vertical/horizontal movement, (which we refer to as cardinal movement) and again in the state of diagonal movement.

Figure 2.3: A jumping maze with red, circled numbers that change the direction of movement.

To handle this new complication, we must modify the model or the algorithm. Modifying the algorithm would involve adding conditional statements to the BFS to allow vertices to be visited twice, once cardinally and once diagonally. However, we prefer to modify the graph instead of the algorithm, for several reasons. The concept of "modeling" a logic maze can also be viewed as a reduction to a state graph (or a maze without state). This facilitates the use of similar components in the maze design function and enables the consideration of similar traits when rating the difficulty of multiple logic mazes, which is discussed in detail in Chapters 3 and 4. Additionally, there is no need to "reinvent the wheel" by creating a new algorithm, as there are already existing graph algorithms that can be used to solve and study maze characteristics. Hence, modifying the graph is more efficient and attractive compared to the alternative of re-implementating every graph algorithm with small modifications.

Combining the modeling concept with a graph library greatly simplifies implementation, enabling one to use all the conveniences that come with graph libraries. This includes treating the graph as an abstract object along with a bug-free

implementation of numerous common graph algorithms. Moreover, as previously mentioned, we emphasize graph modeling and in our instruction we require students to modify the graph instead of the algorithm. To this end, we encourage students to employ a graph library when solving these logic mazes.

Given that we are modifying the graph and not the BFS, consider the following model for the diagonal jumping maze variant M:

Let G be a directed, unweighted graph. For each cell s on the grid with uncircled number n in M, create two vertices in G, $s1$ and $s2$, with $s1$ representing the cardinal state of s and $s2$ representing the diagonal state of s. Add directed edges from $s1$ to the cardinal vertices of distance n from s considering cardinal movement. Add directed edges from $s2$ to the diagonal vertices of distance n considering diagonal movement. The process is similar for circled numbers, except the outgoing edges connect vertices in the two different movement mazes. In this case, add directed edges from $s1$ to the diagonal vertices of distance n considering diagonal movement. Add directed edges from $s2$ to the cardinal vertices of distance n considering cardinal movement.

In the grid, there will be one cell without a number; this is the goal cell. In this variant, the goal cell is always the bottom right corner (though this is flexible and may not be the case in other variants). Create only one vertex to represent the goal. When it is possible to reach the goal in one move whether in the cardinal state or the diagonal state, add an edge to this single goal vertex.

To remove the notion of "state", which is the movement type, we have expanded the number of vertices in the graph relative to the number of grid cells. There are two states per cell, and thus double the number of vertices in the state graph (minus one for the goal), two for each location on the grid/one for each possible state. In effect, we have created two *distinct* mazes, the cardinal maze and the diagonal maze, with circled numbers serving as connections between the two. This technique is common for logic mazes with state changes. We often increase the number of vertices in the graph relative to the puzzle in order to encode its complexity. This expansion requires a constant amount of additional

space per cell, which is not true for all logic mazes; in some cases, the expansion requires a polynomial additional space instead of a constant (see later sections on step-change (Alice) mazes or multi-player mazes).

After constructing the graph, a BFS will find (one of) the shortest path(s) from the start to the goal vertex (the start is always the upper left cell in the cardinal state in this variant). A sample graph is shown in Fig. 2.4, where the vertices are named (state, row, column) where 1 indicates a cardinal state and 2 indicates a diagonal state. The colors and shapes of the vertices will be explained in a later chapter.

2.2.3 Arrow Mazes and Variants

Another common logic maze is the arrow maze. Similar to jumping mazes, the maze is constructed on a grid with the start as the upper left hand cell and the finish as the lower right cell. Instead of numbers on the cells; however, there are arrows. There are eight possible orientations of an arrow, and they are usually denoted using the eight principal directions, with an arrow pointing "North" indicating an upwards arrow, and the other directions can be inferred from this.

The movement in arrow mazes is determined by the direction of the arrow in each cell. As seen in Fig. 2.5, the top left arrow points downwards, thus it is possible to move from the top left cell to any other cell in the first column.

Arrow mazes are modeled in a similar manner to jumping mazes. Let G be a directed, unweighted graph. For each cell s on the grid of arrow maze M, add a vertex representing this cell and directed edges from this vertex to the vertices representing the cells pointed at by the arrow in s. This reduces the arrow maze to a state graph, and we can run a BFS from the start to find a shortest path from the start to the finish.

Abbott presents several variants of the basic arrow maze, including Apollo and Diana and Apollo's Revenge (Abbott, 1990, 16-17). Apollo and Diana introduces a coloring scheme to the basic arrow maze, in which the arrows are colored red and blue, and the solver must alternate between red and blue arrows while moving through the maze (see
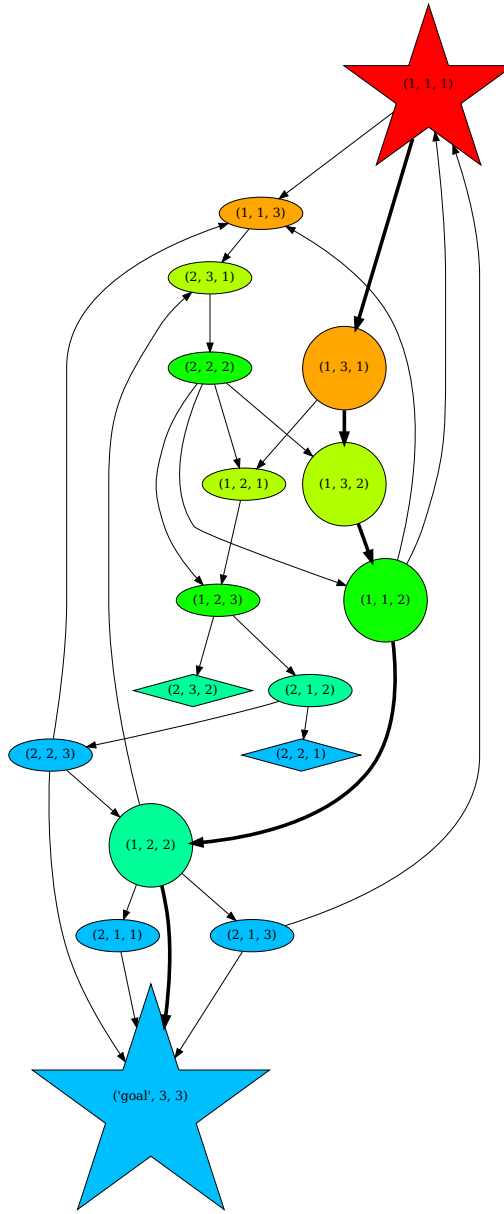
Figure 2.4: Graph model of the jumping maze instance shown in Fig. 2.3.

Fig. 2.6). Adding a coloring scheme of any kind (red-blue, red-green-blue, etc.) does not change the model significantly—when iterating over the cells pointed at by the current arrow, an edge is only added from the current vertex to the vertex representing the child if the color is valid.

Apollo's Revenge introduces a state change similar to the circled numbers in jumping mazes (see Fig. 2.6), where circled arrows change the direction of movement from forward to backward. When a player lands on a circled arrow while moving forward, they start

Figure 2.5: A simple arrow maze.

Figure 2.6: The left shows a colored arrow maze. Solvers must alternate colors when moving through this maze. The right instance depicts a colored arrow maze with movement state changes.

moving out of the *tails* of the arrows, that is, backward, until reaching another circled arrow, at which they resume forward movement. This variant combines a coloring rule with a state-change.

To model a state-changing arrow maze, we can follow a similar approach as for a state-changing jumping maze. Specifically, using the graph model described earlier for jumping mazes with state changes, create a subgraph for each type of movement; the circled arrows connect the two subgraphs to form the overall state graph for the maze.

Abbott introduces what he considers the best of his arrow mazes in *SuperMazes*, the sequel to *Mad Mazes*, and he coins this variant "Arrow Hockey." In an Arrow Hockey maze, the rules of a basic arrow maze apply, except that there is a dime on one of the cells. The goal of the maze is to bump the dime from cell to cell until moving it onto the goal cell. Movement rules are identical to a basic arrow maze (one can move to any cell that the current cell's arrow points at) except that one may not travel over the dime. Although one cannot travel over the dime, the solver can bump the dime by ending a move on the same cell as the dime. In this case, the dime is bumped one cell forward in the direction the solver came from when landing on the cell with the dime. For instance, if the dime was at cell $(3, 3)$ and the player was at $(3, 1)$ with a rightward pointing arrow, moving onto cell $(3, 3)$ would bump the dime to cell $(3, 4)$.

The location of the dime changes the rules of the maze. Therefore, each possible location of the dime represents a different maze. To entirely remove state and reduce this maze to a state graph, we will have n subgraphs, one for every possible location of the dime (which is assumed to be every cell in the maze). This means the graph representation of an arrow hockey instance will have n^2 vertices, where n is the number of cells in the original maze. If we consider a $n \times n$ grid version of the arrow hockey maze, it will have n^4 vertices. This reduction expands the number of vertices and edges, and hence the memory consumption by a polynomial factor based on the size of the original maze, unlike the constant factor observed for previous state-changing mazes. Due to the size of the state graph, arrow hockey mazes with only a thousand cells can become difficult to work with.

Mazes like arrow hockey raise the question of whether it is better to modify the maze-solving algorithm to simulate the maze directly instead of creating a state graph. In contrast to direction state-changing mazes, the separate subgraphs in an arrow hockey instance are very similar, differing only in the location of the dime and the cells pointing at it. Consequently, it may be unnecessary to perform the state graph reduction, which results in consuming a much larger amount of (mostly wasted) space.

Figure 2.7: Connections maze with two colors and four line types.

2.2.4 Connections Mazes

Another maze often used in the undergraduate algorithms class is Grandpa's Transit Map (Abbott, 1990, 18-19), also known as a "Connections Maze" or an "Either-Or maze" (Abbott, 1997, 28-29). Connections mazes resemble a graph and have vertices (villages) and edges (transit lines). The objective is to travel from the start (Startsburg) to the goal (Endenville). From the starting village, one can choose any outgoing transit line (in Abbott's connections mazes, the start only has a single line connecting it to another village, so this is a simple extension of his rules) to take to a connecting vertex, but from that point on, the line on which one chooses to exit a circle must have the same type or the same color (a "free transfer" in Grandpa's Transit Map) as the line used to enter the village. Additionally, this maze has a no-U-turn rule, meaning that one cannot leave a village on the same line used to enter it, even though it (trivially) has the same color and/or type as the line used to enter the village.

In the connections maze depicted in Fig. 2.7, the goal is to move from the starting village (A) to the ending village (G). It is possible to move from the starting vertex (A) to either B or C. If one chooses to move to C from A, then the only option (given the rules above) is to take the red arrow line to B, then to D, then the path splits and one could choose to move to either E or F.

In the model for connections mazes, the transit lines become vertices (two vertices per

13

line in the original maze) and the villages become edges. Consider an arbitrary connections maze M. Let G be a directed, unweighted graph. For every line L that connects villages $v1$ and $v2$ in M, create two vertices in G, one which represents "boarding" transit line L from village $v1$ and another which represents boarding transit line L from village $v2$. For the vertex that represents boarding L from $v1$, add directed edges to the vertices that represent boarding transit lines at $v2$ with the same color or the same type as L, except the vertex that represents boarding L at $v2$ because of the no-U-turn rule. For the vertex that represents boarding L from $v2$, add directed edges to the vertices that represent boarding transit lines at $v1$ with the same color or type as L, except the vertex that represents boarding L at $v1$ because of the no-U-turn rule.

Because it is possible for multiple transit lines to connect to the starting village A, add a unique starting vertex and connect it to the vertices representing boarding any transit line at village A. Similarly, add a unique finish vertex and add edges from the vertices representing the transit lines boarded at all villages that connect to the finish village to this new finish vertex. The state graph representation of the connections maze given in Fig. 2.7 is shown in Fig. 2.8.

Traveling one direction along a transit line is completely different than traveling the opposite direction on the same transit line (these are separate vertices in the state graph representation). This property can be exploited to design clever and challenging mazes that require retracing many of the same transit lines in reverse order.

2.2.5 Multiplayer Mazes and Variants

Many logic mazes involve multiple players, which Abbott calls "pointers." Some examples from Abbott's work include Spacewreck, Meteor Storm, and Theseus and the Minotaur (Abbott, 1990, 32-35). In the undergraduate algorithms class, we use Spacewreck, which has two players, but the model for two players generalizes to mazes with n players.

Spacewreck involves a graph-looking maze with rooms (vertices) and corridors (edges),
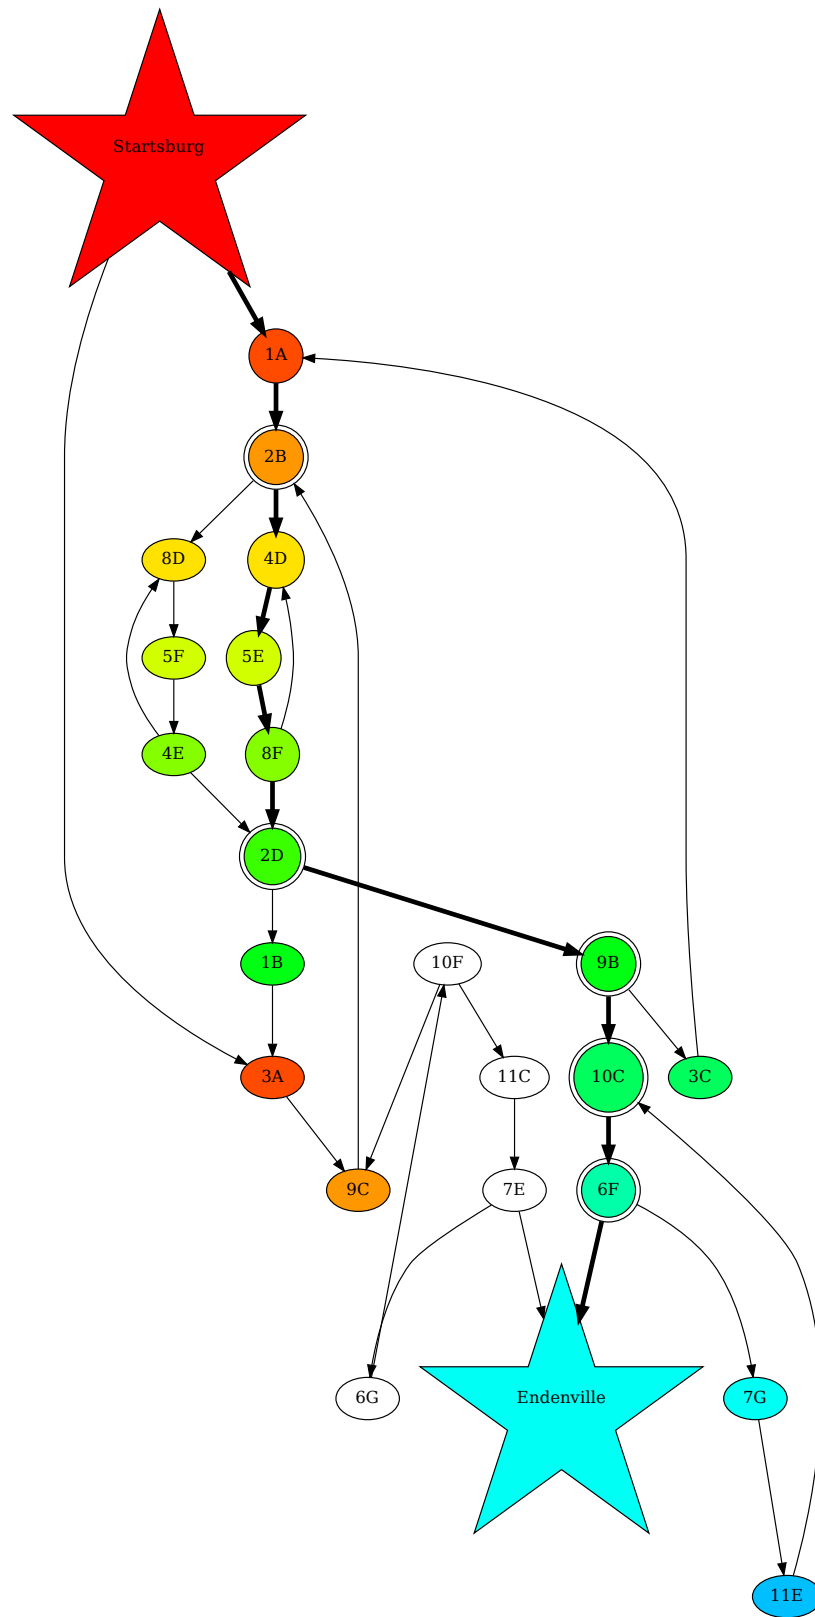
Figure 2.8: State graph model of Fig. 2.7

Figure 2.9: An example Spacewreck maze.

each with specific associated colors. The corridors can only be traversed in one direction in this variant, and one of the two players must reach the room marked "Goal" to complete the maze. Player 1 starts in room A and player 2 starts in room B. Only one of the two players can move at a time. If player 1 is in a room with color c, then player 2 can move through an adjacent corridor of color c to a new room. Similarly, If player 2 is in a room with color c, then player 1 can move through an adjacent corridor of color c to a new room.

Consider the instance shown in Fig. 2.9 where players 1 and 2 start in rooms A (red) and B (blue), respectively. Both players can move, and there is no requirement for them to alternate moves. Player 1 can move down the blue corridor to D because player 2 is in a blue room, or player 2 can move down the red corridor to A because player 1 is in a red room. However, this move results in both players being in room A, which is a dead end because A has no outgoing red corridors. A good notation system to keep track of the maze state is the locations of the two players in alphabetical order. Since the players do not have to alternate and only one of them must reach the goal, it is unnecessary to track

which player is where. The two players start in position AB. Player 1 moves to D, resulting in position BD; after which player 2 moves to A, leaving the players in position AD. $AB \rightarrow BD \rightarrow AD \rightarrow AE \rightarrow AC \rightarrow CD \rightarrow DD \rightarrow DE \rightarrow EE \rightarrow EC \rightarrow GOAL$ is one possible path to reach the goal.

The state graph model for multiplayer mazes is polynomial in the number of players. All possible locations of each player/pointer must be included as vertices, with edges representing the movement of one player. In the case of Spacewreck, the notation system introduced previously is used to denote each vertex - AA, AB, AC ... AF BB BC BD ... FF. We give the two letters that represent each player's location in lexicographic order to avoid repetition of vertices. The vertices are generated by iterating over all pairs of rooms in the original maze in lexicographic order.

In the rules of this multiplayer maze, only one of the players must reach the goal vertex. Therefore, when generating the state graph, any vertex that has one of the two players in the goal room is the same state, which is denoted "GOAL GOAL." For example, "A GOAL", "B GOAL", would be one vertex in the state graph model, "GOAL GOAL."

Let G be a directed, unweighted graph with rooms R. Iterate over the rooms in pairwise fashion to generate all vertices AA, AB, etc. For each vertex, check if the color of room one matches any outgoing corridors from room two. If a match is found, add an edge from the current vertex to the vertex that represents the player moving down that corridor. The same process is repeated for room two and outgoing corridors from room one. If both players are in the same room, then the same edges could be added multiple times, but this graph representation simply ignores duplicate edges (it is not a Multigraph). If the outgoing corridor leads to the GOAL room, then add an edge from the current vertex to "GOAL GOAL," assuming the colors match properly. See Fig. 2.10 for the graph representation of the Spacewreck instance given in Fig. 2.9.

This model is polynomial in the number of players. It has on the order of n^p vertices where p is the number of players. If there was a Spacewreck variant with three players, then the model would contain vertices AAA, AAB, AAC, etc. This maze has no "turns",
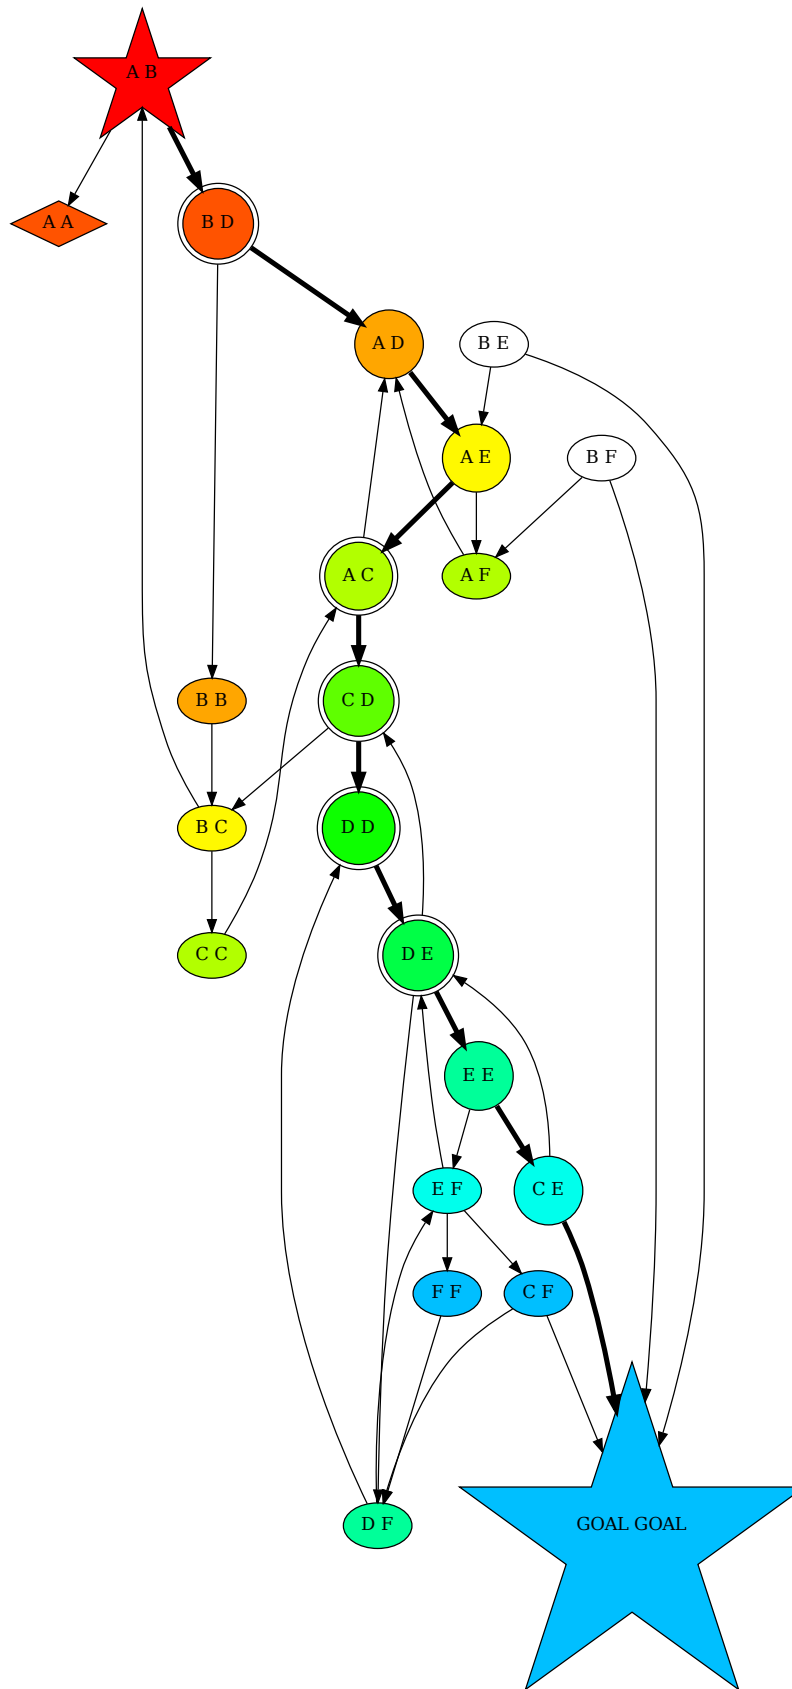
Figure 2.10: State graph model of the maze in Fig. 2.9

i.e., any player can move at any time instead of alternating. If there were turns; however, this would add a constant factor onto the number of vertices in the state $p * n^p$.

2.2.6 Misc. Other Mad Mazes

All logic mazes (that we have found) can be modeled as a state graph and solved with a BFS. Theseus and the Minotaur (Abbott, 1990, 34-35) involves a model similar to Spacewreck, with the additional complexity of determining which player's turn it is at each location pair. Alice in Mazeland (Abbott, 1990, 28-29) is played on a grid similar to arrow and jumping mazes, except the arrows point in multiple directions that indicate the allowed directions of movement. The player's step length starts as one, but certain squares increase the step length by one and others decrease it by one. This model is similar to the state-changing Apollo's Revenge and Jumping Jim mazes—create a separate subgraph for each possible step length, and connect the subgraphs with edges to and from the vertices that change the step length. This model is polynomial in the side length of the grid because the step length can be anything from $1..s - 1$ where s is the maximum side length of the grid. Dice mazes (or Rolling-cube mazes as they are often called) are also played on a grid and have 24 vertices in the graph model per square: six different numbers can be on top of the die, and for each of those six numbers, four different numbers can be facing the bottom (or equivalently, the top) of the grid (Abbott, 1990, 37).

To summarize, all mazes that we have encountered so far can be represented as a state graph using polynomial additional space compared to simulating the maze directly. While it is theoretically possible that a maze exists that requires exponential space and time to perform this reduction, we have not found one yet. It is worth noting that Rolling-cube mazes in which one must visit a subset of squares exactly once have been proven NP-Complete (Buchin et al., 2007, 8-9), but this is fundamentally a different problem than the mazes we have considered, as we attempt to find *any* path from the start to the goal in our puzzles rather than a Hamiltonian path, which is the case in Buchin et al. (2007).

CHAPTER 3

MAZE CHARACTERISTICS

3.1 State Graph Characteristics

The primary objective of this research was to create logic maze instances of varying sizes, both small enough to serve as instructional examples for students to become familiar with the rules of the maze, and large enough to test their implementations.

The state graph reduction applied to all mazes in this work allows for analysis of both the underlying graph representation and the high-level problem instance, as they are reduced to an unweighted directed graph. This unique approach enables the use of an identical objective function when calculating the score of state graphs, regardless of maze type, which is a significant contribution of this work. Additional qualities can then be applied to the higher-level problem instances individually.

A local search implementation resulted in the highest-quality mazes. This is discussed in detail in Chapter 4, but the main purpose of this chapter is to explain and highlight important maze qualities that can be incorporated into an objective function for local search.

Neller et al. (2011) discuss some desired maze attributes that will be introduced, along with additional characteristics identified in our original work. Abbott also outlines several metrics by which he designs his mazes, which will be mentioned as well. We found some additional maze qualities to be challenging based on personal experience, but these have not been verified with human testing. To our knowledge, none of the mazes presented have been used in a study to determine their true difficulty, instead we have selected metrics from Abbott, Neller et al. (2011), and personal experience.

Before delving into specific maze characteristics, it is essential to note when designing a maze for humans to solve, a common solving method is to work backward from the goal instead of forward from the start (Abbott, 1999, 38). Therefore, it follows that the maze ought to be equivalently difficult when attempting to solve it forward or backward,

otherwise it will be easily solved through back-tracing. Considering not only the state graph but also its transpose is critical when evaluating the following maze characteristics. A practical method to enforce the same difficulty in the transpose is to compute the metric value for for both the state graph and its transpose, and add the minimum (maximum) value to the objective function, assuming maximizing (minimizing) the objective function is the goal. For the remainder of this chapter, it is assumed maximizing the objective function is the goal.

Fig. 2.8 is the state graph of Fig. 2.7 based on the reduction given in Chapter 2. An explanation of the colors and shapes in the state graph diagram is given below.

- Shapes
    - The start and goal vertices are star-shaped.
    - Diamond-shaped vertices are unreaching, meaning that it is impossible to reach the finish from them.
    - Vertices on the shortest path (only one path if multiple shortest paths) are circular shaped.
    - Required vertices are double-circled. This means that any path from the start to the finish must pass through these vertices.
    - Edges between vertices on the shortest path are bolded for convenience.
- Color
    - Vertices are shaded according to the distance from the start vertex.
    - Vertices with a higher degree of red coloration indicate shorter distances from the start.
    - Vertices colored similar to the goal vertex are at about the same distance from the start as the goal vertex.
    - Vertices colored a darker blue than the goal are farther away from the start than the goal vertex.
    - Uncolored vertices are unreachable from the start.

3.1.1 Paths, Shortest Paths, and Number of Vertices

The complexity of a maze can be indicated by the number of vertices in its graph representation. Mazes with more vertices tend to be more challenging due to larger traps (explained in a later section), longer true and false paths (the true path leads to the solution while a false path does not), and increased decision points. However, simply adding more vertices without consideration for other aspects of the maze does not necessarily make it more difficult. Therefore, the number of vertices should not be directly included in the objective function, but can be used to adjust the weight of other terms in the function to accurately reflect the size of the maze.

Abbott suggests leaving a substantial amount of space for several long false paths when deciding what portion of the vertices should be involved in the shortest path (Abbott, 1997, 27). From experience with our program, a good number is 15-35%, but exceeding 35% could limit the potential for other characteristics that increase maze difficulty. If a majority of the maze is involved in the shortest path, then over 50% of vertices are "correct" and cannot contribute to other maze traits. To account for this in the objective function, a deduction is applied to the score if the shortest path involves less than 15% or more than 35% of the vertices.

In addition to the length of the shortest path, the existence of multiple shortest paths can affect the solver's motivation. Neller et al. (2011) find that there is a level of satisfaction achieved when one discovers the shortest (or best) solution, and the existence of a unique shortest solution can motivate solvers to continue working on a maze even after solving it (Neller et al., 2011, 193). Therefore, our objective function includes a bonus if there is one shortest path (with some exceptions, see specific maze qualities below).

3.1.2 Branching

Branching is a relatively intuitive characteristic that refers to the number of different locations one could be in after exactly X moves from the start. It is similar to the idea of a "branching factor" in the growth of a search tree/space, and is used to ensure there are no

sections of forced moves at the beginning or end of the maze, at which point the maze might as well start after the forced moves. It also helps to confuse solvers working both forward and backward if one could be in 14 different possible locations after only three moves.

When calculating the branching score, it is necessary to determine, how many possible moves to take into account. An effective method to calculate a branching score is to adjust the number of moves considered from the start position based on the length of the shortest path. The longer the shortest path, the higher the values of X when calculating the total number of possible locations after X moves from the start. 10-15% of the shortest path length is a good starting point.

An additional aspect to take into account when calculating branching is if repeats will be considered (e.g. there is a vertex that can be reached after either 2 or 3 moves from the starting vertex). Typically, vertices should not be double counted when calculating branching because a human maze solver who is less than 10% into a reasonably sized maze will likely notice they are revisiting a state they could have reached from the start in fewer moves, and thus choose an alternate path. While it is possible to consider repeats, branching scores that include them tend to disproportionately impact the overall maze score.

3.1.3   Reachability

A *reaching* vertex v is a vertex from which it is possible to reach the goal g. A *reachable* vertex is a vertex that can be reached from the start s (Neller et al., 2011, 192). A traversal from the start (finish) on the state graph (transpose) can be used to calculate the portion of reachable (reaching) vertices. Sixty percent is a good minimum portion of the maze to be reachable (reaching) from the start (finish). This restriction can greatly influence the creation of traps, as traps are often only accessible when moving in one direction through the maze.

Reachability also measures the efficiency of a maze because the number of vertices in

the state graph is a function of the puzzle size, which usually does not change while generating an instance. Often, one chooses a size (e.g., 30 rooms and 72 corridors in a Spacewreck maze) and the generator does not modify the size, it merely changes the other properties of the maze (colors, which rooms the corridors connect, etc.). Using the formula for combinations with replacement with $n = r - 1$ and $k = 2$, and then adding one for the goal, a state graph of a Spacewreck maze with r rooms will have $\binom{r}{2} + 1$ vertices. A 30 room instance will have 436 vertices. If only 200 of these vertices were reachable and/or reaching, this would be a poor use of resources.

Vertices that are both unreachable and unreaching are denoted *isolates*, and these represent wasted resources—vertices in the state graph that cannot be reached from the start nor backward from the finish. The puzzle instance ought to be redesigned so these are included as a part of the maze. Isolates negatively influence the score of the maze in objective functions as a result.

If a logic maze allows bidirectional movement and the state graph is undirected (similar to a classic maze), all reachable vertices will also be reaching (even if one must pass through the start to reach them in the opposite direction or vice versa). In this case, there will only be isolates and vertices that are both reachable and reaching.

3.1.4 Traps, Holes, and Whirlpools

From the definitions of reachable and reaching vertices, we can define several different traps to confuse and hinder the maze solver. A dead end of a maze is a set of one, or many, reachable, unreaching vertices. A reverse dead end is a set of reaching, unreachable vertices in the state graph.

A *black hole* is a set of strongly connected, reachable, unreaching vertices in the state graph. In effect, it is a false path that ends in loop(s) instead of at a singular dead end. Black holes, especially large black holes, can significantly increase a maze's difficulty because a maze solver may spend a lot of time in the trap before realizing there is no escape. This definition is slightly different than the one given in Neller et al. (2011)

because we wanted to focus on the strongly connected vertices, which is the core of the trap, and neglect the fringes. Black holes usually force the solver to restart the maze once they realize there is no path to the solution because the solver does not remember how they initially entered the black hole (Neller et al., 2011, 193).

A *white hole* is a set of strongly connected, reaching, unreachable vertices in the state graph. It is identical to a black hole when considering the transpose of the state graph (with the start becoming the finish and vice versa). Only maze solvers working backward from the solution can find themselves in white holes. Again, these increase the difficulty of a maze and help confuse backward-minded maze solvers, but not to the extent of black holes. Fig. 2.8 has a small white hole: the vertices 10F, 11C, 7E, and 6G make up a four vertex white hole.

In general, black holes make a maze more difficult than white holes because if the maze cannot be easily solved in a brute-force manner by moving forward or backward, most maze solvers will use analytical/systematic techniques to solve the maze. These techniques are almost always applied to the maze when solving forward from the start rather than backward from the finish.

These types of traps only affect one direction of solving the maze. Reverse dead ends and white holes have no impact on individuals solving the in the forward direction because they cannot be reached when moving forward. Similarly, dead ends and black holes do not impede solvers attempting to move backward from the solution. These traps directly conflict with the reachability score, as their existence implies that there are vertices that cannot be reached when moving in both directions. However, the state graphs of many of these puzzles, especially if the puzzles are reasonably sized, contain enough vertices (100+ vertices in the state graph) that the increase in difficulty from the traps is worth sacrificing reachability in one direction for a subset of vertices.

A *whirlpool* is a set of strongly connected, reaching, reachable vertices. In effect, it is a hole that can be reached in both directions, and the placement of such traps is more difficult and important. A whirlpool located near to the start is vulnerable to back-tracing

from the finish, and the same can be said of a whirlpool close to the finish. This is because a solver can often see a path to the goal state when within a few moves of it, even if they have to pass through a large trap to reach it.

When assessing maze difficulty, it is preferred to include some type of hole instead of a dead end. Holes are inherently more troublesome than a simple dead end that halts all progress—one often wastes considerable time moving around a hole and must restart the maze. In contrast, at simple dead ends one can usually retrace their steps before that decision was made.

It is important to avoid overly large black/white holes that make one direction of movement much easier than the other. Typically, reachability constraints prevent a maze from having a single large hole that can be easily solved through back-tracing and vice versa, and applying the method discussed above (taking the minimum of the black and white hole scores before adding onto the objective function) often results in a balanced maze with sizeable black and white holes, or neither.

In a logic maze with bidirectional movement, such as Meteor Storm (Abbott, 1999, 33), there cannot be black holes and white holes. Instead, the only type of trap that can exist are whirlpools.

3.1.5   Decisions, Required Vertices, Bridges/Dominance

The presence of traps such as black/white holes and whirlpools in a maze does not guarantee it is difficult to solve. We have generated many mazes where multiple traps exist, but a solver may not encounter them unless they are unlucky while a lucky solver may never encounter these traps.

We want to maximize the chances for a solver to become lost in traps for them to truly increase the maze's difficulty. The first solution that comes to mind is to consider the *decisions* that a maze solver must make along the shortest path. Consider each vertex on the shortest path n. From n, there are several cases for each immediate descendent d (the resulting vertex of each outgoing edge of n):

- d is the next vertex on the shortest path.

- d is in an unreaching trap (black hole/dead end).

- d is in a reaching trap (whirlpool) or on a suboptimal path to the solution.

- d is a previous vertex on the shortest path.

We aim to have as many decision points as possible that lead into reaching/unreaching traps or suboptimal paths. However, there are several issues with merely counting the number of outgoing edges from each vertex on the shortest path that satisfy this requirement. The simplest problem would be (depending on the maze) the decisions within a few vertices of the solution vertex are often irrelevant because it is clear by inspection what decisions to make when one is only a few moves away from the solution. This can vary depending on the maze. For instance, in the connections mazes shown previously, it is easy to see the path to the solution when one is within 3-4 vertices of the goal, but in jumping mazes it is not as clear and one may need to be within 1-2 vertices of the solution to see the path. Not considering decisions for vertices within three moves of the solution is a good rule of thumb to use in this context. When creating an objective function that considers decisions, it is up to the designer to choose the move cutoff. However, the mazes presented in this work are at a level of difficulty such that most if not all solvers will attempt back-tracing, which means that they will be at least somewhat familiar with the paths of length 3-5 working backward from the finish.

Another consideration is that not all decision points should be weighted equally. If one edge leads into a five vertex black hole and another into a 30 vertex black hole, the latter should have a more significant impact on the score. Various method can be used to address this issue which will be discussed further in the next section.

Decisions are one of the primary factors in determining the difficulty of a maze (Neller et al., 2011, 194). To ensure that decisions factored into the score of a maze are relevant to all solvers, it is important to consider only the decisions that every solver, regardless of path chosen, must make. However, if there are multiple shortest paths or multiple paths to the solution, not all solvers may have to consider every decision, which could falsely

increase the perceived difficulty of a maze. Therefore, we need to determine which vertices are "required" vertices R that are present on all paths from the start vertex s to the goal vertex g. Removing any vertex in R would result in no path existing from s to g.

This question has been studied and solved in theory related to control-flow graphs, and is referred to as *Dominance*. A vertex v *dominates* another vertex u if v lies on every path from the entry vertex to u.

We need to determine which vertices dominate the goal vertex with an entry node of the starting vertex. Let D denote the vertices that dominate the goal with an entry point of the start vertex. If a vertex $n \in D$ (n dominates the goal vertex), then any dominator d of n also dominates the goal vertex, that is, $d \in D$. Thus, we can construct a *dominator tree* of all vertices that dominate the goal. The immediate dominator is the parent of a vertex in the dominator tree. We calculate the dominators of the goal vertex by calculating the immediate dominator of the goal vertex id_{goal} and then the immediate dominator of id_{goal}, etc., until reaching the start vertex.

Dominance has been used to show the safety of code-reordering operations and can be used in control flow graphs to determine which lines of code must execute before others. Cooper et al. (2006) give a $\mathcal{O}(V^2)$ algorithm that in practice runs faster on graphs with less than 1000 vertices than the classical $\mathcal{O}(E \log(V))$ Lengauer-Tarjan Algorithm (Cooper et al., 2006, 1-2). State graphs of typical mazes do not usually exceed this number of vertices. After computing the dominator tree, we can identify the vertices that dominate the goal vertex. By only including the decision scores associated with these vertices, we can ensure the mazes perceived difficulty is based on decisions that all solvers must consider.

3.1.6 Long False Paths (Decisions continued)

Abbott directly mentions leaving space for what he denotes "long false paths" repeatedly in his work on designing mazes (Abbott, 1997, 26-27), (Abbott, 1990, 36-40). A false path is a path that does not lead to the goal. Holes, whirlpools, and forks from the solution path that waste time are all considered false paths.

Long false paths correspond well to suboptimal decisions at the required vertices. To score these decisions, it is important to consider the degree of incorrectness. When a suboptimal decision is made, the resulting vertex d can fall into one of three categories:

1. d is a previous vertex on the shortest path.

2. d is in an unreaching trap (black hole/dead end).

3. d is in a reaching trap (whirlpool) or on a suboptimal path to the solution.

The score associated with the first case should be minimal, as many solvers will remember their previous 10-20 moves, so a decision that leads backward on the correct path to an already-encountered vertex does not make the maze particularly more difficult. Even if the shortest path is of length 80 and the 50th vertex has an edge to the 5th vertex, a solver is likely to recognize the vertex is close to the start and not choose this edge.

In the second and third cases, the score associated with the decision is based on the amount of time a solver spends moving through the trap without making progress. For an unreaching trap, the score is relative to how long it takes the solver to realize they are repeating the same path and need to restart. The algorithmic approach is to determine the maximum distance a solver can move in the trap without retracing their steps. Once the solver starts retracing their steps, they will realize they are stuck and need to find a way out, probably by restarting the maze.

In the third case, the solver has entered a reaching trap, and the score is determined by calculating the furthest distance the solver can move without getting closer to the solution than the vertex at which they made the suboptimal decision, and without retracing their steps further back along the shortest path. For example, consider a solver at a vertex on the shortest path that is 20 moves away from the goal. They make a suboptimal decision and move to a vertex that is 40 moves away from the goal. The maximum score that could be added as a result of this decision is $40 - 20 = 20$. Alternatively, if the new vertex leads into a trap without retracing the solver's steps, we could potentially add the trap score here as well.

Based on the discussion in the previous section, we will apply these decision scoring

criteria to only the required vertices, that is, the vertices that dominate the goal vertex, and we will not include required vertices that are three moves away (or less) from the goal vertex.

3.2 Mad Maze Instance Characteristics

Mazes often have characteristics unique to the problem instance itself that ought to be included in the objective function. Because the state graph abstracts away some qualities of the maze instance, these characteristics cannot be captured by the state graph and must be separately included in the objective function of each specific maze type.

3.2.1 Color

An even distribution of colors is often desired in logic mazes that have multiple colors, such as Connections or Spacewreck mazes. For instance, if a Connections or Spacewreck maze had four colors, it would be undesirable for most of the rooms/villages and corridors/transit lines to be the same color. To ensure an even distribution, the number of appearances of each color can be counted, and the difference between the minimum and maximum number of appearances deducted from the score, and this difference can be squared or cubed if necessary. However, in Spacewreck mazes, an even distribution of colors may lead to additional isolates because it limits possible moves. Thus, the designer must decide which is more important.

3.2.2 State Changes

To increase the difficulty of mazes that involve state changes, such as jumping and arrow mazes, it is important to maximize the number of state changes that occur along the shortest path (or perhaps maximize the number of state changes between successive vertices in the dominator tree). It would be undesirable to not require the solver to change the movement direction between cardinal/diagonal and forward/backward. Thus, we can count the number of state changes along the shortest path/required vertices, and add the product of this quantity and a multiplier (5-10 is usually sufficient) to determine the additional score that results from a state change.

Figure 3.1: A jumping maze generated by a function that maximizes state changes without limiting the number of circled cells.

However, this can result in mazes with a large majority of the grid squares being circled (all of them in some cases), which is suboptimal as shown in Fig. 3.1. To encourage both state changes along the required vertices and that an optimal portion of the grid squares become locations of state changes, we choose a desirable portion of the grid squares (such as 20%) to be state changes. Then calculate the difference between the actual number of state-change grid squares and the desired portion. Square (or cube) this difference and subtract it from the score in the objective function.

3.2.3 Doubling Back and U-turns

In all of the mad mazes presented in Chapter 2, it is possible to revisit the same location in the maze multiple times, and this has been previously explained for jumping, arrow, and connections mazes. In Spacewreck mazes, because state is determined by the locations of both Lucky and Rocky, one player can revisit the same location (such as A) multiple times in an entirely separate section of the state graph. This can disorient and confuse solvers because it feels like they are retracing their steps even though they are not. In the Spacewreck maze instance shown in Fig. 3.2, Abbott designed the maze so the solver must trace around one initial loop (A-J-O-N-H-D-A) two times with one pointer,

another loop (B-G-K-F) *three* times with the other pointer, and only one exit from these dual loops leads to the solution (Abbott, 1990, 37).

The concept of "state" that logic mazes introduce enable this re-visitation, or doubling back, which Abbott uses in many of his mazes to confuse the solver (Abbott, 1990, 36). In the case of Spacewreck, we can track the visited vertices and count the number of times the players revisit them. In connections mazes, we can keep track of the visited villages and count the number of times a village is revisited. An especially tricky aspect of connections mazes is that traveling one direction on a transit line is an entirely separate vertex in the state graph than traveling the other direction on the line. This idea enables one to create entertaining mazes where one traces a large path out into the maze, makes a U-turn in some fashion, follows the same exact transit lines back to near the starting village, and then finds a short path to the finish that is now possible. Abbott uses this concept in his Grandpa's Transit Map instance. The solver begins the maze at Startsburg on the red curvy line, and can transfer to either the blue curvy line to the west or the red straight line south at the next village intersection (indicated with a black arrow in Fig. 3.3). However, any path to Endenville (the goal) *must* return to this village intersection and take the blue straight line to the east, the one line that cannot be accessed when entering the maze from Startsburg. Therefore, the solver must initially move out into the maze, make a u-turn, and retrace their steps back to this village, where they can board the blue straight-line transit line and continue eastward towards the goal.

One caveat of this maze structure is that there will always be two possible shortest paths—the transit lines that one uses to make the u-turn can be traversed in both directions. For this type of maze, it is recommended to allow for two or less shortest paths so these u-turn type instances can be generated.

In jumping mazes, specifically with cardinal/diagonal state changes, it can be especially desirable to double back to the start square along the shortest path (or as a required vertex) but in the diagonal state instead of the cardinal state. Typically, a modest increase to the score equal to the number of squares in the grid is a good amount to reward for this
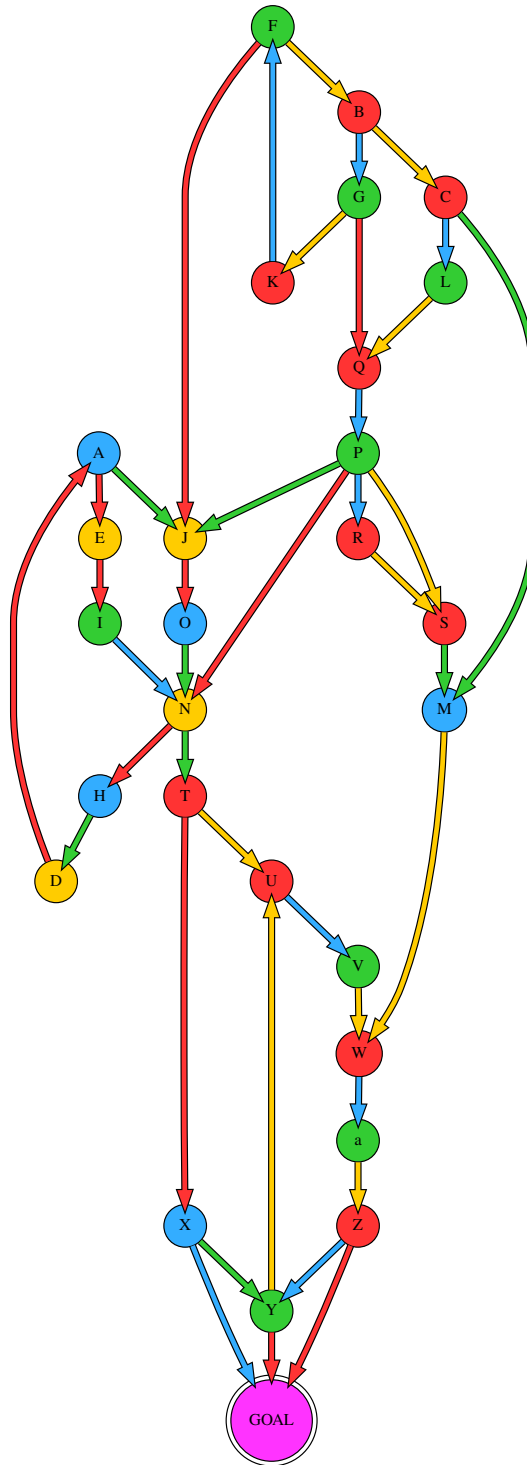
Figure 3.2: A maze depicting the Spacewreck instance given in Mad Mazes (Abbott, 1990, 32).

Figure 3.3: The top section of Abbott's Grandpa's Transit Map connections maze (Abbott, 1990, 19). The arrow points to the village mentioned when discussing transit lines exited from the start.

occurring. However, this cannot occur in arrow mazes because the starting square arrow must point onto the maze, causing its tail to point off the maze.

The concept of u-turns is similar to doubling back, except that it is more unique to the grid mazes discussed (arrow and jumping mazes) because it involves a disorientation of the solver not from revisiting the same square in multiple states, but rather multiple repeated movements in the same row, column, or diagonal. Abbott uses this concept in many of his grid mazes. In the Apollo's Revenge Maze, the shortest path moves up and down a column, reverses direction, and then moves up and down the same column again (Abbott, 1990, 36). Given the shortest path (or adjacent required vertices), it is possible to calculate the amount of moves that require u-turns, moving either up and down the same column repeatedly, or along the same diagonal. The jumping maze given in Fig. 3.4 makes use of many u-turns.

The maze given in Fig. 3.4 begins at (1,1). The path to the goal then moves south to (7,1), northeast to (6,2), east to (6,6), and then doubles back for the next 6 moves along the main diagonal: (6,6) to (2,2) to (5,5) to (1,1) to (7,7) to (4,4) to (6,6). These u-turns along the main diagonal will confuse many solvers, and it also brings a degree of surprise and enjoyment to solving the maze not present in other mazes.

6	1	5	(4)	(4)	2	1	2
1	3	2	2	(1)	6	(3)	5
2	5	4	4	(2)	2	(5)	6
2	6	(1)	2	3	2	1	3
3	2	2	4	4	4	2	4
2	(4)	2	4	3	(4)	1	2
(1)	6	4	2	3	2	3	4
(7)	2	(2)	5	1	1	(1)	GOAL

Figure 3.4: Jumping maze with an abundant quantity of u-turns.

3.3 Component Graph

Generating the component graph or *condensation* of the state graph becomes useful as mazes grow larger and it becomes difficult to identify the state graph characteristics by inspection. The component graph is created by contracting each strongly connected component into a single vertex. Edge labels are then added to show the number of inter-edges between the components in the original graph, and self-loop edges are added to show the number of intra-edges in each component: edges that are between vertices *within* each component. The same coloring and shape rules that were applied to the complete state graph are also applied to the component graph.

In the first maze given in Fig. 3.5, the component graph reveals two small strongly connected components, or two smaller mini-mazes, with a four vertex white hole. For the second maze, the component graph does not tell us a whole lot, other than the maze solver starts within an 18-vertex strongly connected component with 29 intra-edges and two possible ways out (edges) that lead to the solution vertex.

While these component graphs are simple, they are for smaller maze instances. Larger mazes such as the jumping maze in Fig. 3.4 or the Spacewreck instance in Fig. 3.2 have

Figure 3.5: Component graphs of the connections and spacewreck mazes given in Figures 2.7, 2.8, 2.9, 2.10.

component graphs that are too unwieldy and large to be shown in this paper. Removing all single-vertex strongly connected components from the condensation that are unreachable *or* unreaching and do not have both incoming and outgoing edges leads to component graphs that can be shown in this paper. The simplified component graph for Fig. 3.4 (the u-turn jumping maze) is provided in Fig. 3.6.

In Fig. 3.6, the maze solver starts the puzzle in a 40 vertex strongly connected component with 68 intra-edges, with exactly one edge that leads out of this whirlpool to the goal vertex. There is a 30-vertex black hole with 15 direct entrances (edges) from the starting scc and additional entrances that pass through other unreaching sccs. There is also a 25-vertex white hole with five direct entrances and around four times that number of indirect entrances through other unreachable sccs. The major challenging in solving this
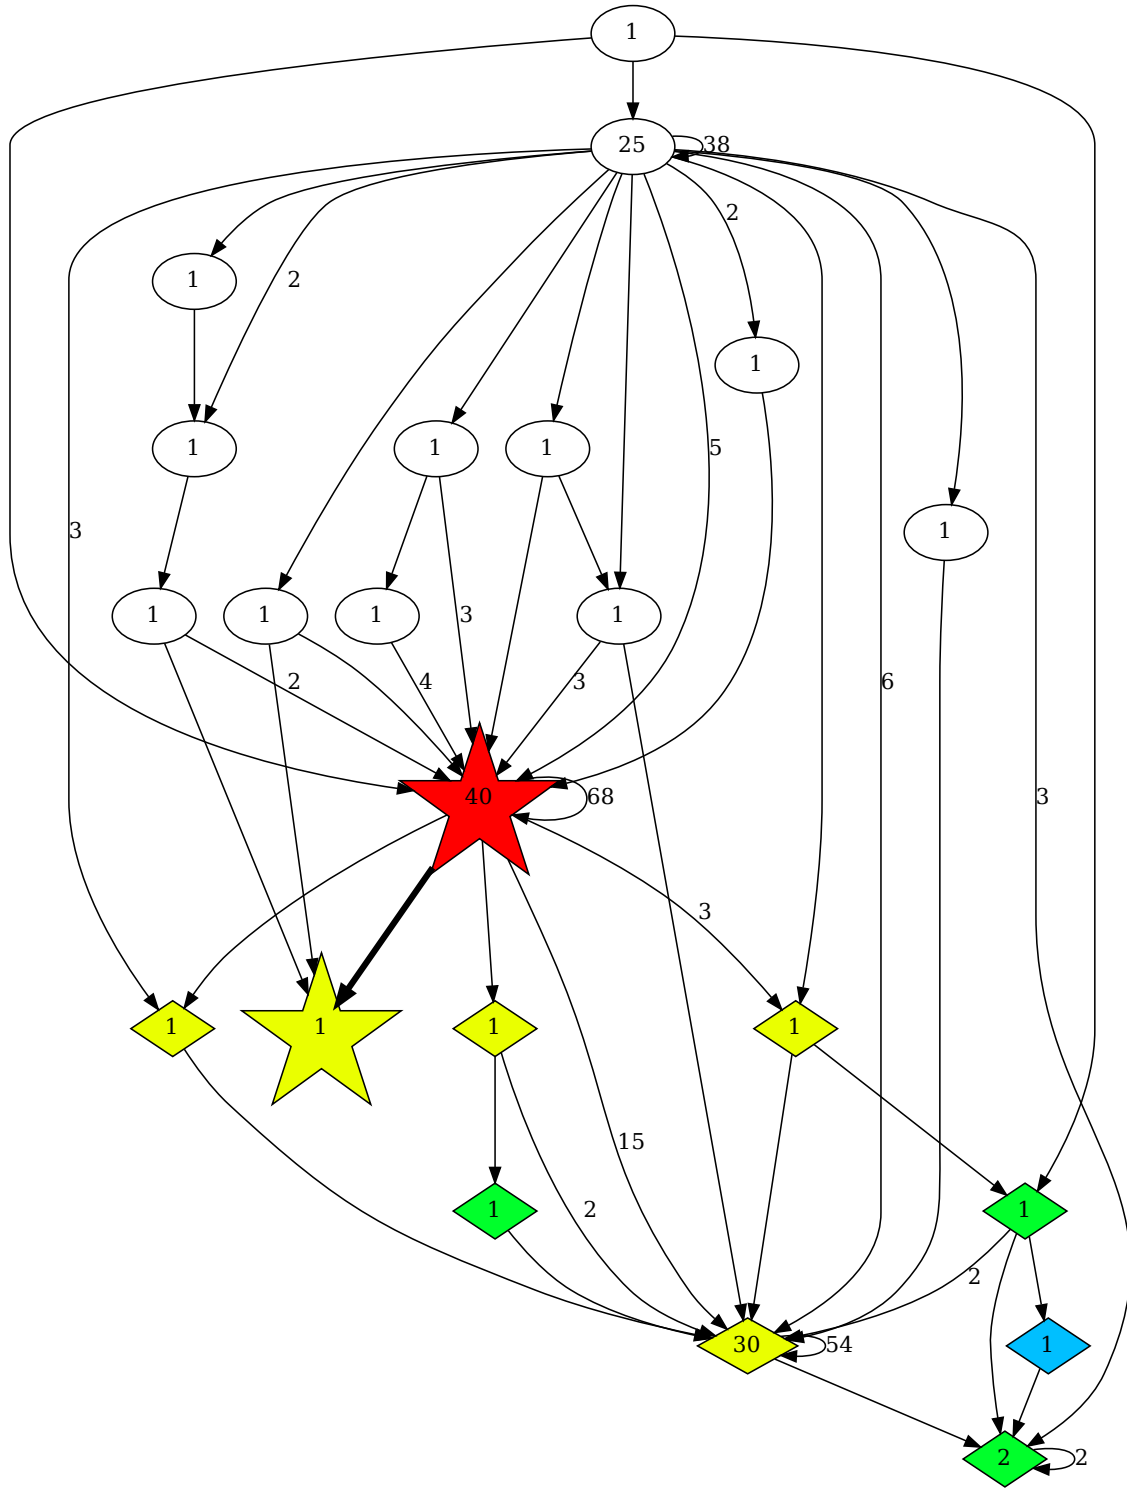
Figure 3.6: Component graph of Fig. 3.4 with single-vertex unreachable or unreaching sccs removed that did not have both incoming and outgoing edges.

maze arises from the sole edge that connects the starting scc to the goal and the presence of two signficant holes.

CHAPTER 4

IMPLEMENTATION AND SEARCH

4.1 Implementation

The implementation of mad mazes is well-suited for an object-oriented approach, which allows for the sharing of common characteristics, such as being played on a grid or color constraints. The root object contains the state graph logic because all mazes have a state graph representation. More specific types of mazes can either inherit from this object or add components that they use, such as a grid, colors, etc. See Fig. 4.1 for a UML representation of the architecture that we used for this implementation.
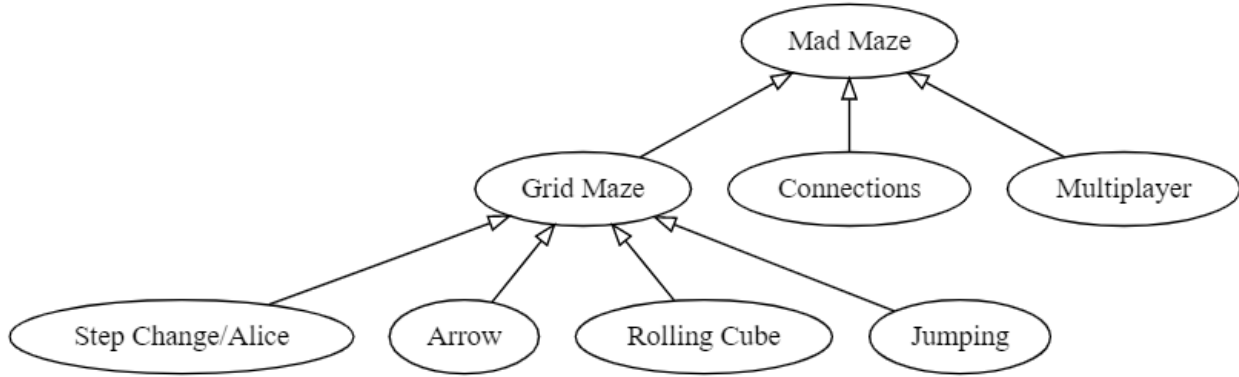
Figure 4.1: Potential inheritance scheme

The root logic maze object contains the state graph objective function while the objects that pertain to each specific maze puzzle call their parent's objective function to rate the state graph, and themselves have operations to consider the attributes of the higher-level puzzle instance (such as an even distribution of colors). The root also contains the state graph visualization and comparison functionality, state graph analysis, and abstract methods for reading input, outputting solutions and generation that the children implement.

4.1.1 Program Structure

One efficient way to implement the detection of attributes discussed in Chapter 3 is to use the state graph reduction combined with a robust graph library such as Networkx in Python 3. This library includes pre-implemented versions of all the necessary functions to detect and score state graph attributes, such as dominance, traversals, and distance calculations. The nodes in Networkx can be any object with user-defined attributes. The library also includes algorithms for traversals, immediate dominators, shortest path lengths, flow, descendants, ancestors, bfs/dfs trees, and more. Additionally, Networkx provides visualization capabilities, allowing for graphs to be exported to Graphviz for visualization purposes. For more information on Networkx, refer to its documentation (Hagberg et al., 2008, 11).

Functions that operate on mazes to rate attributes (beyond the graph library functions) should be implemented as helper functions that are given the maze instance and the maze start and finish. This enables reuse of the same function when considering both the forward and transpose of the state graph by passing in the transpose, the finish as the start, and vice versa.

4.2 Fully Random Generation

A purely random generator was attempted for each logic maze. We chose the size of the maze and randomly generated the characteristics of each location/edge in the maze (grid number, arrow direction, color, type). Overall, this method produced unsatisfying results. In most instances generated in this manner, there was no path from the start to the finish, and even when a solution path existed, it was often trivial, and large portions of the maze were unreachable and unreaching. We omit details on this generation type because the results were unfruitful.

4.3 Intelligent Random Generation

These methods of generation do have random components, but are designed to be more intelligent than the fully random generation—rather than generating an entire puzzle

instance from scratch, we focus on improving certain aspects of a maze.

4.3.1 Automation of manual generation

Abbott discusses several methods of maze generation in both *Mad Mazes* and *SuperMazes*. Besides focusing on the traits (long false paths, decisions, etc.) introduced in Chapter 3, Abbott gives more direct instructions on generating a rolling cube maze.

He suggests first laying out the "true" path, or the solution path through the maze from start to finish. This true path should not take up the majority of the maze to leave room for false paths; Abbott directly mentions that a maze "won't be interesting" unless it has several complex false paths. After completing the true path, Abbott recommends laying out several long false paths. Finally, the maze is completed by choosing numbers for the remaining squares that create the "most and longest false paths" possible (Abbott, 1997, 27).

The first attempt at "intelligent" random maze generation was to automate algorithms such as the one introduced in the previous paragraph, and we first tried this for 8×8 jumping mazes with a diagonal state change.

Inputs: Minimum shortest path length, Start, Finish, Probability of circling.

Output: A jumping maze instance with circled squares for diagonal state changes.

1. True path generation: The truth path was generated using a backtracking search. A randomized traversal was conducted from the starting point to generate the true path through the puzzle.

    a) When choosing numbers for squares in jumping puzzles, the numbers must be sufficiently small to allow movement in ideally both the cardinal and diagonal states.

    b) The random numbers that ensure movement is possible are iterated over, then all reachable square locations (children) after one move are calculated.

    c) These children are randomized then iterated over.

    d) We randomly determine if this square will be circled, then loop over all possible

numbers (in a random order) that ensure there are reachable locations from this square in both cardinal and diagonal movement.

e) If the number we have chosen to put onto a child would create a new shortest path to the parent or a path to the solution that is shorter than the minimum desired length, then that number is skipped. If all numbers would cause this to happen, then we change the child to circled (assuming it was uncircled) and vice versa, then attempt all possible numbers again. If it is still impossible to add a number to this square, then we backtrack to the parent and try adding a different child to the solution path. If we cannot add any of the current children to the solution path, then we try a different number for the parent.

We have a backtracking solution for generation of the shortest path. The worst case complexity of this algorithm is certainly exponential and it would never finish in some instances, but in practice it quickly terminates.

2. False paths

a) To attempt to create false paths from the start and the finish as Abbott mentions, we calculate all squares that can be reached from the starting point but are not on the shortest path and do not yet have numbers. Then, we attempt to number these squares and then the subsequent children in such a manner that the shortest path is not intersected. We do this both moving forwards from the starting point of the maze and backwards from the finish of the maze, considering distances of five from the start/finish.

After creating these false paths, we attempt to number the remaining squares using the possible number and circle combinations that would not create a new shortest path to the finish.

While this algorithm resulted in mazes that were certainly improved from the randomly generated mazes, the inconsistent run times and variability in the difficulty of results was undesirable. It was at this point that the authors found the paper by Neller et al. (2011)

that recommended the use of local search with an objective function that highlighted maze qualities, and we found this method to be much more successful.

4.3.2   Local Search

As mentioned previously, this type of maze generation does not start from scratch but rather takes a currently generated maze instance and makes small modifications that increase the value of an objective function. When defining local search solutions to problems, there are generally three items to define:

1. An objective function that returns a score given a solution instance to the problem. We seek to either minimize or maximize this function.

2. A notion of a neighborhood, or a set of small modifications made to a given solution to turn it into another solution then rated by the objective function. Typically, a solution has multiple neighbors.

3. A search algorithm, that is, a method to choose between the neighbors of a given solution.

In the context of maze generation, the objective function will be as defined in Chapter 3 and contain metrics to rate both the state graph and the maze instance. The neighborhood definition turns one maze instance into another, which involves making a small change to the maze that is simultaneously reflected in the state graph. We considered generating an ideal state graph and turning it back into a maze, but it is difficult and challenging to determine the minimum requirements for a state graph to be able to map it to a maze instance.

Consider a jumping maze with diagonal state changes. If we take an arbitrary state graph and try to determine if it is a jumping maze state graph with diagonal state changes, there are several requirements that we must have. In effect, we will determine if it is possible to map the state graph onto a grid.

1. We must be able to partition the state graph into two equal-sized partitions (except for the solution/goal, which should have no outgoing edges) such that all edges are

within each partition except for the edges outgoing from what would be considered the "circled" nodes, which must all point into the other partition.

2. Let p1 denote the cardinal partition and p2 denote the diagonal partition.

3. All outgoing edges from each node must travel the same "distance" when mapped onto the grid, so that a number can be placed onto the grid square.

4. Each vertex in p1 must have an equivalent vertex in p2 (the cardinal and diagonal versions) that has outgoing edges that represent diagonal movement with the proper distance.

5. Circled grid squares must have edges from p1 travel diagonally into p2 and edges from p2 travel cardinally into p1.

Attempting to generate the state graph without taking the puzzle rules into account and then trying to map it onto the maze is likely to fail due to the strict conditions. Hence, a better approach is to consider the puzzle rules and define neighbor states based on the puzzle rather than the state graph.

4.3.3 Search Options

There is a plethora of search algorithms that one can use to select neighbor states, all with some measure of success. We focused on two primary search methods: stochastic local search (with random restart) and simulated annealing. Stochastic local search (sls) involves iterating over ones neighbors in a random order, and accepting the first neighbor encountered with a higher score or rating (according to the objective function) than the current state. If all neighbors result in the same or a worse state, the algorithm terminates. We used a fully random algorithm to generate a maze instance, then applied the sls search algorithm, and took the best result. Also, starting the search from the maze instances provided in *Mad Mazes* and *SuperMazes* proved successful as well. Sls often results in being stuck in local maxima/plateaus in the search space because it cannot, at any point in the algorithm, accept a worse state than the current, unlike simulated annealing (Russell and Norvig, 2022, 111-114).

Simulated annealing (sa) is similar to sls in that it involves iterating over one's neighbor's in a random order and always accepts better neighbors, but it also may accept inferior neighbors. The probability of acceptance is based on how poor the neighbor state is (a lot worse is less likely to be accepted) combined with a "temperature." The probability of acceptance is given by the following formula $P(A) = e^{\Delta S/T}$ where ΔS represents the difference in value between the neighbor state and the current state and T is the temperature. When ΔS is negative, that is, the neighbor state is worse than the current state, at high temperatures there will be a high probability of accepting this worse state. However, at low temperatures, it is unlikely that this worse neighbor would be accepted. As $T \to 0$, the probability of accepting a worse state approaches zero (Russell and Norvig, 2022, 114-115).

Typically, the temperature starts at a maximum value then decreases as the algorithm runs (with each iteration), so near the start of the algorithm it behaves similar to a random walk before starting to climb to higher states. We implemented a decreasing temperature schedule based on the number of iterations. For the first 500 iterations the temperature is 100, then decreases by half every thousand iterations until reaching 5000, at which point it is set to 0.01 (so close to zero that the algorithm behaves like a stochastic local search and only accepts higher or same-rated states).

The temperature is very flexible. It should not be set too high to avoid simulating a fully random walk because that would defeat the purpose of starting from an already excellent maze, but outside of this constraint, as long as the temperature decreases over time, the algorithm finds good results.

With both sa and sls, we run the search until it terminates: sls when all neighbor states are worse than the current, and sa when the temperature has become zero, the algorithm is essentially sls and terminates in the same fashion.

Another search option is to generate all neighbor states and accept the one that leads to the highest score, but the overhead in computing all neighbor states is very high and this algorithm was often too slow in practice to yield useful results, especially as puzzle

instance size increased.

4.3.4 Neighbor States

We define the neighbor states for each type of maze introduced in Chapter 2 below. These are implemented in Python using generator functions, so we iterate over all possible neighbor states as they are generated one by one; thus, when we choose to accept a state as a successor we have not done the work of generating *all* neighbor states, only those that were evaluated before accepting a successor.

1. Jumping Mazes

    For jumping mazes, a neighbor state is simply changing the number and/or the circling of a particular square on the grid. We iterate over the squares in a random order, followed by the numbers and the two circling options. There is an 80% chance to try not circling before circling to help avoid an excess of circled grid squares.

    Once a square is chosen, a new number and circling option are selected, and the outgoing edges from each vertex in the state graph corresponding to that square are removed and stored. Then, we calculate the edges that would result from the new square, number, and circling option and add them to the state graph. Next, we run the objective function to score the new state graph and maze instance. Depending on the search function being used, we may accept or reject this neighbor state. If we accept it, we move to the next iteration immediately. However, if we reject it, we remove the new edges added to the state graph, re-add the old edges, and reset the grid square back to its original number and circling.

2. Arrow Mazes

    Arrow maze generation is nearly identical to jumping mazes. A random square, arrow direction, color (if present), and circling option are selected (i.e. iterated over in a random order within the neighbor generator function). The state graph and maze are updated with the change, the new instance scored with the objective function, and if accepted we move to the next iteration, otherwise all changes are reset before moving onto the next neighbor.

3. Connections Mazes

   Connections maze generation is more complex. In order for the mazes to viewable and solvable by humans, a grid structure is enforced in the maze instance (not the state graph) so that transit lines are only between adjacent vertices in this grid structure (up to 8 adjacencies for each location).

   Because the state graph is defined by the transit lines and not the villages themselves, we select a random transit line, color, and line type. Then we choose which village to keep the transit line connected to (one of the two villages to which it is currently connected), and randomly choose a new successor village from the grid structure. Note that two villages can only be connected by at most one transit line. We store the original characteristics and make the modifications to the state graph and puzzle instance before using the objective function to obtain a new score. This score is passed to the search function which determines if we accept or reject the new state. Accepting leads to another iteration of search, rejecting leads to resetting all changes before generating the next neighbor.

4. Multiplayer/Spacewreck/Meteor Storm Mazes

   Similar to connections mazes, a grid structure is enforced in the puzzle instance so that the generated maze is clean and easily viewable by humans. We randomly choose a corridor and a corridor color, then randomly choose one of the two rooms to keep this corridor connected to (call this room R), randomly choose a new room from the adjacencies of R that this corridor will now connect to. We also choose new colors for these two rooms. Make the changes in the puzzle instance and state graph, get a new score from the objective function, determine if we accept or reject, reset for next neighbor or move onto the next iteration as mentioned in the other three mazes.

We discard neighbor states that result in the same current puzzle instance and state graph.

4.4 Example Objective Function

In this section, we will provide an example search function used for state graph generation that focuses on decisions. Let n denote the number of vertices in the state graph. It is best (if possible) to try and scale terms using n to provide weightings that represent importance. Note that we will be attempting to maximize the score of the maze (and not minimize).

The score is initialized to zero. We define the following function to score the state graph:

1. If there does not exist a path from the start to the finish, then subtract n^3 from the score.

2. Otherwise, count the number of shortest paths from the start to the finish. If there is only one, add $10 * n$ to the score. If there are multiple, do not add $10 * n$.

3. If the length of the shortest path(s) is greater than $n * 0.35$ or less than $n * 0.15$, subtract n^2 from the score. This ensures that the path does not consume too much of the state graph so we can use other parts of it for false paths, and also that we do not have a trivial solution path.

4. For the following characteristics, they must be calculated for both the state graph and its transpose to ensure adequate difficulty when moving forwards from the start and backwards from the finish. Thus, calculate both cases and take the minimum of the two, and add that to the score.

5. Next we consider the branching and reachability scores. Using a BFS, calculate all descendant vertices and categorize them by distance from the starting vertex. Then, calculate the sum of vertices cr that are reachable after a certain number of moves, (10-20% of the shortest path length is a good starting point). Add cr to the score (after calculating cr_t and taking the minimum of cr and cr_t, where cr_t is the same statistic but calculated on the transpose with the finish as the starting point).

   For reachability, count the number of vertices r reachable from the start. As mentioned previously, we calculate the proportion reachable from the finish in the

48

transpose graph. Call this r_t. Add the minimum of r and r_t to the score. Then, union the results, which gives all vertices that are reachable or reaching. Add the number of vertices that are either reachable or reaching to the score. Determine all vertices that are in the state graph that are neither reachable nor reaching. Subtract this quantity squared from the score—these represent wasted vertices (isolates).

6. The most important factor in determining the difficulty of a maze, given all other traits are equal, is the decision score that determines exactly what decisions every maze solver must make. Compute the required nodes that every path from the start to the finish must go through using the dominance algorithm discussed in Chapter 3 and given in Cooper et al. (2006). Then, evaluate the decisions at these vertices in both the state graph and its transpose, as discussed in Chapter 3, take the minimum of the two, multiply by two, and add to the score.

7. Now we consider dead ends. A dead end is a reaching or reachable vertex in a state graph that does not have any outgoing edges, (or equivalently in the transpose, does not have any incoming edges). Holes make better traps than dead ends, as mentioned in Chapter 3. Subtract the number of dead ends times five from the score.

This objective function generates state graphs that focus on maximizing the decision score. The higher-level maze instance terms along with generation results are presented in Chapter 5.

CHAPTER 5

GENERATION RESULTS

Although we use the objective function defined in Chapter 4 to rate the state graphs, we still need to consider the higher-level maze instances when generating, and these additions onto the objective function are defined below for each different type of mad maze, followed by the generation results.

The size of the puzzle instance was chosen for each maze so the state graph has ≈ 128 vertices to enable comparison across the different mazes. This means an 8×8 grid for the grid mazes (jumping and arrow), a connections maze with ≈ 63 transit lines (this will have $63 * 2 = 126$ vertices in the state graph, plus two for the start and end vertices). There will be 32 villages to maintain a similar transit line : village ratio as in Grandpa's Transit Map, which has 36 villages and 70 transit lines. The Spacewreck instances will have 16 rooms, which equates to 121 vertices. Seventeen rooms would result in 137 vertices.

5.1 Jumping Maze - Jumping Jim's Encore

The instance objective function terms are defined as follows:

1. It is undesirable to have large areas of the grid that share the same number (Neller et al., 2011, 194-195). Initialize adj-num and adj-circle to zero. Iterate over the grid and calculate the following: for every grid square, evaluate its (eight) neighbors. For each neighbor that shares the same number, add one to adj-num. If the current node is circled and encounters a neighbor that is also circled, add one to adj-circle. Count the number of circled grid cells as well.

2. Subtract adj-num and adj-circle^2 from the score. Calculate the desirable number of grid squares to be circled (we used 20%). Subtract the difference between the desired number of circled locations and the number of circled locations cubed from the score.

3. Now we consider doubly visiting grid squares, u-turns, and state changes (from cardinal to diagonal to cardinal, etc.). These traits are defined in Chapter 3. We calculate the number of squares that are visited in both the cardinal and diagonal

states, the number of u-turns along the shortest path, and the number of state changes between diagonal and cardinal. Add the number of state changes times five and the sum of u-turns plus doubly visited squares times 10 to the score.

4. Subtract n^2 from the score if the starting vertex is circled. If this is allowed, remove this term.

This function yields excellent jumping puzzle instances and state graphs. Table 5.1 shows the results of several mazes generated using this objective function and simulated annealing with random restart.

Table 5.1 offers several items to consider. The path score (SP score) is the same for each instance. Every maze has a single shortest path, and thus all have the $10 * n$ resulting SP score. We purposefully included the $10 * n$ to enforce this condition, and we appear to have been successful. In effect, we can consider the base score for all of these mazes to be 1270, which is $10 * n$.

The branching score gives a small bonus onto the score as desired to help prevent long sections of forced moves. If one was generating instances where a large branching factor is more important, then one could add a constant to this term, square it, or multiply it by n, depending on its relative importance.

There are less reachable/reaching vertices as a result of holes in the highest rated mazes. This is a common theme we have noticed when generating instances: one can maximize reachability at the cost of traps, which results in a higher reachability score at the cost of a lower decisions score, or one can maximize hole sizes, which results in a higher decision score at the cost of a lower reachability score. In this objective function, we are more heavily weighting the decisions score—it has a constant factor of two applied to it, and the way we are calculating it generally leads to higher numbers than the reachability score, hence, the best mazes sacrifice reachability for traps and a higher decisions score. Compare maze one to maze seven. Maze seven has the highest reachability score (241) of the instances presented in Table 5.1 with the highest proportion of reachable (91.3%), reaching (90.6%), and both reachable and reaching (81.9%) vertices. However, it has the

Table 5.1: Jumping generation maze results are shown below with the primary components of the scores highlighted in **bold**. SP = shortest path, BH = black hole, WH = white hole, R or R = Reachable or Reaching, R and R = Reachable and Reaching. The hole entrances are only to the largest hole. The required decisions is the sum of the required vertices out degrees, which represents the minimum quantity of choices maze solvers will consider. Fwd and bwd decisions are the forward and backward decisions, and the other metrics are previously specified in this chapter and Chapter 2.

Metric/Maze	1	2	3	4	5	6	7	Avg.
SP Length	36	38	34	27	22	36	43	33.7
SP Quantity	1	1	1	1	1	1	1	1
SP Score	**1270**	**1270**	**1270**	**1270**	**1270**	**1270**	**1270**	**1270**
Branching Score	**24**	**31**	**31**	**15**	**15**	**15**	**17**	**21.1**
Reachable (%)	68.5	66.9	72.4	72.4	69.3	88.2	91.3	75.6
Reaching (%)	62.2	66.1	74.0	89.0	92.9	89.8	90.6	80.7
R or R (%)	98.4	99.2	99.2	100	98.4	99.2	100	99.2
R and R (%)	32.3	33.9	47.2	61.4	63.8	78.7	81.9	57.0
Reachability Score	**199**	**208**	**216**	**218**	**208**	**236**	**241**	**218**
Required Vertices	36	38	33	26	20	32	39	32.0
Required Decisions	82	87	78	60	50	70	84	73.0
Largest BH	35	30	22	8	2	2	2	14.4
BH Entrances	18	16	15	2	1	1	1	7.7
Largest WH	25	24	22	4	17	2	2	13.7
WH Entrances	25	23	14	10	18	6	1	13.9
Fwd Decisions	1024	902	834	832	794	558	418	766
Bwd Decisions	1046	990	850	834	794	560	418	784.6
Decisions Score	**1024**	**902**	**834**	**832**	**794**	**558**	**418**	**766.0**
Dead End Score	-50	-50	-60	-50	-80	-60	-60	-58.6
State Graph Score	**2467**	**2361**	**2291**	**2285**	**2207**	**2019**	**1886**	**2216.6**
Circled (%)	23.4	21.9	25	21.9	23.4	21.9	23.4	23.0
State Changes	7	7	10	4	6	12	12	8.3
Double Visited	8	10	6	4	3	8	11	7.1
U-turns	9	9	10	9	5	15	17	10.6
Instance Score	**143**	**186**	**137**	**99**	**54**	**235**	**276**	**161.4**
Overall Score	**2610**	**2547**	**2428**	**2384**	**2261**	**2254**	**2162**	**2378**

lowest decision score.

The decision score dominates the overall score, usually representing about 40%. We claimed previously that decisions are the most important factor when considering the difficulty of a state graph, and this principle is reflected in the scoring criteria we have adopted.

The highest rated mazes generally have a combination of the most required decisions, (which may or may not be the most required vertices), large black and white holes, and many entrances to these traps from required vertices (which is the BH/WH entrances table entry). The top two mazes exhibit these traits, and both have the highest decision scores. A very low R and R score is indicative of mazes with large unreaching/unreachable traps; in the most optimal version of these mazes (according to this objective function), only the shortest path vertices are reachable and reaching, and all other vertices are part of large strongly connected holes.

However, large holes are not required to have a large decision score, as maze four shows. Maze four combines several small white holes (of sizes three, three, and four) with one medium size black hole and a large whirlpool. When suboptimal forward decisions do not lead into the black hole, they lead into a whirlpool that costs the solver a lot of time while suboptimal backward decisions lead into false paths and unreachable traps. The higher R and R score indicates the presence of a large portion of vertices that are reachable and reaching, that is, a whirlpool in this case.

We can see the constraint of subtracting (squared) the difference between the desired (20%) and actual number of circled grid cells has been very successful in constraining the instances to ±5% of the desired value. Adding a bonus for state changes, doubly visited squares, and u-turns has led to their prevalence in the generated instances, with averages of 8.3, 7.1, and 10.6 respectively. The additional bonus for doubly visiting the start square was also successful with nearly every instance doubly visiting the starting square. Adjacent and circled overlap (not shown in the table) have been sufficiently avoided as well.

As designed, the state graph score plays a much larger role than the instance score. When using the same objective function to generate multiple different mazes, it would defeat the purpose of sharing a state graph objective function if it did not play a more significant role than the individual maze instance terms.

The most exceptional of these mazes is depicted in Fig. 5.1 and its component graph is shown in Fig. 5.2. The state graph has the following statistics:

Figure 5.1: Exceptional jumping maze instance.

- $n = 127$ vertices.
- One shortest path of length $36/127 = 29\%$ of the state graph.
- After five moves from the start//finish, one can be in $11//7$ distinct locations not considering repeats.
- $87/127 = 68\%$ of vertices are reachable from the start.
- $79/127 = 62\%$ of vertices are reaching (reachable moving backwards from the finish).
- $125/127 = 98\%$ of vertices are reachable or reaching, so there is only two isolates.
- $41/127 = 32\%$ of vertices are reachable and reaching: there are only five vertices that are not on the shortest path that can be reached in both movement directions, all other vertices are part of holes. This indicates near-maximal hole sizes.
- Every vertex on the shortest path dominates the goal. Thus, shortest path vertices and required vertices are synonymous for this puzzle instance, maximizing the possible decision score.
- The largest black hole is 35 vertices with 18 total entrances.

Figure 5.2: Component graph of the excellent jumping maze instance.

- The largest white hole is 25 vertices with 25 total entrances.

This state graph is exceptional according to the metrics we have defined. It has substantial branching, an optimal number of required vertices, massive black and white hole traps and false paths with many entrances, near-perfect reachability given the large traps, and a near-optimal shortest path length. The higher-level maze instance is excellent as well, with the following statistics:

- 15/64 circled grid cells = 23%.
- There are seven swaps between diagonal and cardinal movement, eight grid squares visited in both diagonal and cardinal states, and nine u-turns.

Almost exactly 20% of the grid squares are circled, which was the percentage chosen when generating this maze in the objective function, and there are few areas with large groups of the same number or circling (compare to Fig. 3.1). The shortest path requires seven state changes on a path of length 36, averaging one every five moves. Eight grid squares are visited both diagonally and cardinally, including the starting square (1, 1), and there are nine u-turns along the shortest path, both of which are outstanding for a path of length 36.

This instance is exceptional and it is unrealistic to expect such outcomes on every generation. The problem is over constrained due to the many conflicting traits, which makes optimization particularly difficult. However, the instances presented in Table 5.1 plus 30 additional instances with an average score of about 2250 were generated within an hour, thus demonstrating this method's consistency.

5.2   Arrow Maze - Apollo's Revenge

The objective function is identical to the jumping maze, including the puzzle instance terms. Because both of these mazes are played on a grid, the puzzle instance metrics optimized in the jumping mazes have a near-identical definition once we consider arrow directions instead of numbers. Circling percentage, state changes, doubling back, and u-turns are identical, and we want to avoid large groups of arrows pointing the same direction and large groups of circled grid cells.

The generation results are shown in Table 5.2. Similar to jumping mazes, we generated many instances and tabulated the statistics for seven, all within the top 25%.

Table 5.2: Apollo's Revenge generation maze results are shown below with the primary components of the scores highlighted in **bold**. SP = shortest path, BH = black hole, WH = white hole, R or R = Reachable or Reaching, R and R = Reachable and Reaching.

Metric/Maze	1	2	3	4	5	6	7	Avg.
SP Length	30	31	31	35	37	32	31	32.4
SP Quantity	1	1	1	1	1	1	1	1
SP Score	**1270**	**1270**	**1270**	**1270**	**1270**	**1270**	**1270**	**1270**
Branching Score	**30**	**32**	**37**	**15**	**26**	**31**	**36**	**29.6**
Reachable (%)	70.1	69.3	70.1	94.5	78.0	90.6	67.7	77.2
Reaching (%)	59.8	63.0	61.4	79.5	77.2	76.4	59.8	68.2
R or R (%)	95.3	98.4	96.9	99.2	96.9	97.6	96.9	97.3
R and R (%)	34.6	33.9	34.6	74.8	58.3	69.3	30.7	48.0
Reachability Score	**160**	**200**	**184**	**225**	**204**	**211**	**182**	**195.1**
Required Vertices	30	31	31	35	36	28	31	31.7
Required Decisions	83	83	88	95	89	68	82	84.0
Largest BH	14	14	14	1	1	1	5	7.1
BH Entrances	14	18	23	0	0	0	15	10.0
Largest WH	24	15	25	1	12	1	19	13.9
WH Entrances	20	24	20	0	13	0	18	13.6
Fwd Decisions	836	718	708	488	494	456	490	598.6
Bwd Decisions	840	734	766	494	494	440	578	620.9
Decisions Score	**836**	**718**	**708**	**488**	**494**	**440**	**490**	**596.3**
Dead End Score	-160	-115	-130	-140	-130	-145	-130	-135.7
State Graph Score	**2136**	**2105**	**2069**	**1858**	**1864**	**1807**	**1848**	**1955.3**
Circled (%)	21.9	23.4	21.9	21.9	20.3	23.4	23.4	22.3
State Changes	10	9	14	13	12	13	14	12.1
Double Visited	3	2	4	5	7	4	4	4.1
U-turns	10	6	4	9	9	9	4	7.3
Instance Score	**105**	**65**	**79**	**146**	**138**	**117**	**66**	**102.3**
Overall Score	**2241**	**2170**	**2148**	**2004**	**2002**	**1924**	**1914**	**2057.6**

Arrow mazes are similar to jumping mazes in most traits, but they generally have lower scores (2058 average vs 2378 average). The ability to pick any descendant the current arrow points at significantly limits the ability to create traps without an escape. From the statistics in Table 5.2, one can see that only the best of the mazes generated have nontrivial black and white holes, which can substantially increase the decision score. As a result, arrow mazes have lower decision scores on average when compared to jumping

57

Figure 5.3: Best Apollo's Revenge maze instance.

mazes, even though there is on average a higher number (about 10 more on average) of required decisions—these decisions do not carry as much weight as they do not usually lead into large traps.

Due to the forward and backward movement, it is impossible to move out of both vertices in the state graph that represent the corners. It is similarly difficult for other squares on the edges of the grid unless they point directly North or South. This leads to a much larger dead ends penalty when compared to jumping mazes (-135.7 average vs -58.6 average). The best instance generated (Fig. 5.3 and Fig. 5.4) follows the paradigm of large holes with many entrances, and thus a poor R and R score. However, it has much fewer holes and more dead ends when compared to the exceptional jumping maze instance, which follows from the properties of the maze.
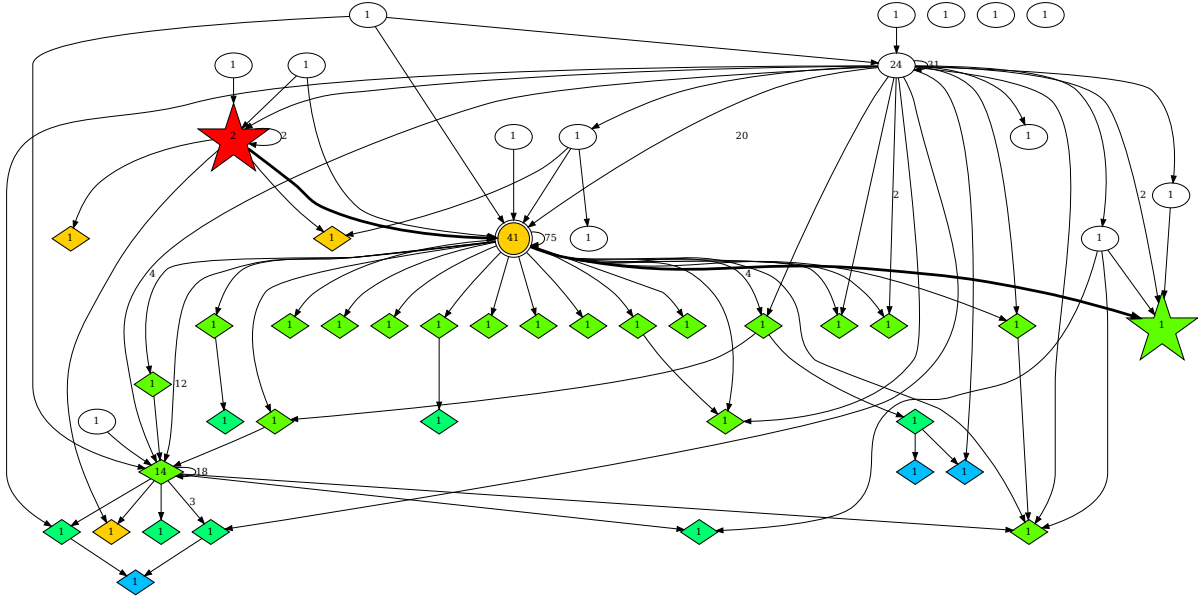
Figure 5.4: Component graph of the best Apollo's Revenge instance. Compared to the jumping maze instance in Fig. 5.2, it has many more dead ends and smaller holes.

5.3 Connections Maze - Grandpa's Transit Map

Three terms are added to the objective function to account for the higher level traits of connections mazes. Firstly, we add $10 * n$ to the score if there are two or less shortest paths. This enables us to generate the interesting mazes as mentioned previously that trace a long path out into the maze, turn around, and follow the same transit lines in reverse order. There will always be two possible ways to traverse this loop, so we allow for two shortest paths. Additionally, we count the number of times we visit the same village multiple times and when we use a transit line in both directions (these account for doubling back in connections mazes). We add three times the number of same village visits and 10 times the number of double line visits to the score.

Lastly, we want to ensure there is an equal distribution of village line types and colors. Count the number of each type/color, calculate the difference between the most prevalent and least prevalent line type, multiply this difference by the number of different line types, and finally subtract this number from the score. Repeat the process for the line and village

59

colors.

A primary feature of connections mazes is connectivity and reachability. It is often the case that nearly 100 percent of the vertices are both reachable and reaching, leading to the maximum possible reachability score. This can conflict with the potential decision score that would result from large traps, however. When generating connections maze instances with this objective function, there is a large group of local maxima that result from maximizing the reachability score, and it can be difficult for the search to instead find an instance with large traps, which leads to a much larger decision and overall score. After generating many mazes using simulated annealing and random restart, we found several mazes with large traps (about 1 in 10 had large traps).

Another possible method is to run simulated annealing with random restart until it creates a maze with some traps, and then running simulated annealing on that instance with a restricted temperature (say, only 5-10), encouraging only small steps down to avoid the large local maxima.

The results for this objective function are given in Table 5.3. All of the mazes have two shortest paths and line doubling, which means they all have some form of loop that can be traced in multiple directions as discussed previously.

The reachability score average is 223, higher than jumping and arrow mazes, indeed every maze has zero isolates and this score is below 210 for only one instance. The bidirectionality inherit within transit lines enables better reachability, which is why we see this general increase in both the reachability statistics and scores. Even the mazes with large holes have better reachability scores and statistics when compared to the jumping and arrow mazes with similar characteristics. Similarly, the dead ends penalty is much lower on average (-22.1 compared to -135.7 and -58) for the same reasons.

We see the same two paradigms of this objective function emerge in the connections maze results as with the grid mazes. The highest-scoring mazes (top four) have large holes with many entrances and required decisions that lead into the holes. The average to medium scoring mazes (five to six) exhibit excellent reachability and few unreaching traps

Table 5.3: Grandpa's Transit generation results are shown below. There are 128 vertices in each state graph, and SP Addback is the $10 * n$ given from having two shortest paths. When a line/village is visited multiple times, this is counted in Village/Line Doubling. The color traits penalize the score for unequal color distributions.

Metric/Maze	1	2	3	4	5	6	7	Avg
SP Length	39	40	38	38	29	30	35	35.6
SP Quantity	2	2	2	2	2	2	2	2
SP Score	**0**	**0**	**0**	**0**	**0**	**0**	**0**	**0**
Branching Score	**26**	**35**	**28**	**31**	**7**	**10**	**30**	**23.9**
Reachable (%)	69.5	70.3	68.0	68.8	88.3	96.9	64.1	75.1
Reaching (%)	69.5	70.3	69.5	70.3	88.3	95.3	69.5	76.1
R or R (%)	100	100	100	100	100	100	100	100
R and R (%)	39.1	40.6	37.5	39.1	76.6	92.2	33.6	51.2
Reachability Score	**216**	**217**	**214**	**215**	**240**	**249**	**209**	**222.9**
Required Vertices	36	37	35	35	26	27	32	32.6
Required Decisions	92	91	90	88	51	53	71	76.6
Largest BH	29	29	29	28	4	3	30	21.7
BH Entrances	35	32	36	31	5	1	28	24
Largest WH	29	29	29	28	4	1	30	21.4
WH Entrances	35	31	37	32	5	2	29	24.4
Fwd Decisions	1922	1776	1768	1522	1362	1148	1228	1532.3
Bwd Decisions	1916	1760	1814	1602	1546	1148	1166	1564.6
Decisions Score	**1916**	**1760**	**1768**	**1522**	**1362**	**1148**	**1166**	**1520.3**
Dead End Score	-25	-25	-15	-20	-15	-20	-35	-22.1
State Graph Score	**2133**	**1987**	**1995**	**1748**	**1594**	**1387**	**1370**	**1744.9**
Village Doubling	21	21	19	19	12	12	14	16.9
Line Doubling	14	14	13	13	9	8	10	11.6
Village Colors	-32	-12	-12	-4	-8	-4	-36	-15.4
Line Colors	-15	0	-6	-6	0	0	-9	-5.1
Instance Score	**156**	**191**	**169**	**177**	**118**	**112**	**97**	**145.7**
SP Addback	**1280**	**1280**	**1280**	**1280**	**1280**	**1280**	**1280**	**1280**
Overall Score	**3569**	**3458**	**3444**	**3205**	**2992**	**2779**	**2747**	**3170.6**

although some have whirlpools.

The highest-scoring maze generated by this algorithm is nearly optimal when considering its characteristics. The state graph in Fig. 5.6 has sufficient branching to avoid forced moves at the start (we can see that one edge leads directly into the black hole, so there is a 50% chance of a solver immediately entering it), two shortest paths (so it gets the $10 * n$ SP Addback), zero isolates, and still a substantial reachability score (216). The largest contributor to the score is the decision score, and we can see why this is the case

Figure 5.5: Granpa's Transit Map instance one. Note that E is an isolated village without any transit lines. This means we could potentially increase the number of transit lines when generating these instances. One of the two shortest paths is A F L Q M L R N S R W X S O T S Y d c X d Y S T O S X W R S N R L M Q W b f.

from the state graph. Of the length 39 shortest path, 36 are required vertices. At these 36 vertices, every incorrect decision leads either into the black hole or requires many steps to be retraced in the whirlpool, and when retracing these steps one can easily end up in the black hole as well.

The decision score is not the only exceptional trait: this maze also exhibits the tricky property mentioned previously: the shortest path retraces 14 of the same lines in both directions. It traces out a path into the maze, turns around, and then traces the same lines

in the reverse order before turning towards the solution.

5.4   Multiplayer Maze - Spacewreck

The additional objective function terms for Spacewreck are similar to Grandpa's Transit Map. We start by counting the number of rooms and corridors of each color. However, unlike the previous maze, we must strictly enforce an even color distribution or the search will result in 90% of the rooms and corridors being the same color. One can understand intuitively based on the rules of this maze that a nearly monochromatic maze will have the highest reachability and the fewest isolates. However, this often leads to long sequences of forced (or obvious) moves. Consider the case where both players are in rooms of the most prevalent color (which happens often), and the outgoing corridor(s) from p1 are not the most prevalent color. P2 must find a path to a less prevalent-colored room following edges of the most prevalent color. This often leads to long sequences of forced and obvious moves.

However, enforcing the even color distribution too strictly can result in a large number of isolates, which is wasted space in the state graph as mentioned previously. Thus, calculate the difference between the least prevalent color and the most prevalent color for both rooms and corridors separately. Subtract each of these values **cubed** from the score to enforce equal coloring. Cubing this term instead of multiply by n allows there to be small differences in color distribution without a large penalty, but the penalty grows rapidly as the differences increase. This results in mazes with color differences of 1-4, which is ideal.

Similar to doubling back in previous mazes, we count the number of times that a repeated room appears in the shortest path and add this quantity times three to the score. A tricky aspect of this maze is it is possible for both players to be in the same room, which may seem counterintuitive at first to human solvers. We count the number of times both players are in the same room and add this number times 20 to the score.

We used the state graph objective function modified with these additional instance terms to generate many Spacewreck instances. The statistics of seven instances are presented in Table 5.4.

Table 5.4: Spacewreck generation maze results are shown below. Num Colors is the number of colors used in creation of the maze instance (two or three). Room repeating is the count of the same room being visited multiple times, Same Room is the count of the number of times the two players must be in the same room, and the colors are penalties for uneven color distributions.

Metric/Maze	1	2	3	4	5	6	7	Avg.
Num Corridors	42	42	42	42	42	31	42	40.4
Num Colors	2	3	3	2	3	2	2	2.4
SP Length	26	22	27	23	31	31	25	26.4
SP Quantity	1	1	1	1	1	1	1	1
SP Score	**1210**	**1210**	**1210**	**1210**	**1210**	**1210**	**1210**	**1210**
Branching Score	**28**	**13**	**25**	**21**	**24**	**25**	**17**	**21.9**
Reachable (%)	56.2	59.5	63.6	71.9	63.6	73.6	86.8	67.9
Reaching (%)	71.1	57.9	62.0	73.6	65.3	81.8	93.4	72.1
R or R (%)	99.2	96.7	97.5	100.0	95.9	99.2	97.5	98.0
R and R (%)	28.1	20.7	28.1	45.5	33.1	56.2	82.6	42.0
Reachability Score	**186**	**170**	**183**	**207**	**167**	**207**	**213**	**190.4**
Required Vertices	26	22	27	21	31	31	24	26.0
Required Decisions	78	48	57	61	68	73	68	64.7
Largest BH	16	28	28	21	20	14	3	18.6
BH Entrances	29	20	18	8	21	11	2	15.6
Largest WH	23	27	26	8	23	12	3	17.4
WH Entrances	31	20	25	17	24	11	5	19.0
Fwd Decisions	814	866	854	778	646	546	506	715.7
Bwd Decisions	918	858	886	772	740	618	504	756.6
Decisions Score	**814**	**858**	**854**	**772**	**646**	**546**	**504**	**713.4**
Dead End Score	-55	-70	-105	-45	-85	-30	-50	-62.9
State Graph Score	**2183**	**2181**	**2167**	**2165**	**1962**	**1958**	**1894**	**2072.9**
Room Doubling	11	10	10	10	13	14	10	11.1
Same Room	2	3	1	1	5	1	2	2.1
Room Colors	-1	-8	0	-27	0	-1	-1	-5.4
Corridor Colors	0	-8	-8	-8	0	-1	0	-3.6
Instance Score	**72**	**74**	**42**	**15**	**139**	**60**	**69**	**67.3**
Overall Score	**2255**	**2255**	**2209**	**2180**	**2101**	**2018**	**1963**	**2140.1**

The results are similar to the other mazes with the two paradigms (reachability vs holes) although most of the mazes fall into the holes category. Compared to the other maze types, the scores are generally lower, similar to arrow mazes. These mazes are slightly handicapped by having seven fewer vertices in the state graph, which would be an increase of 70 on the SP Score and lead to increases in other categories as well.

The top two mazes are tied. It is interesting that the first maze has a lower decisions score, but makes up the difference with a combination of reachability and branching. Maze three also has a larger decision score than maze one, but maze one makes up this difference through less dead ends and a higher instance score. Maze one demonstrates that even the lower magnitude point metrics can significantly impact the score.

There is a difference in the number of corridors and also the number of colors in many of these instances, and yet they all scored well. Two mazes tied for the highest score of 2,255, one with two colors and one with three colors. Choosing the number of colors for the rooms and corridors was challenging. Abbott uses four colors for 28 rooms and 40 corridors, but his Spacewreck instance has many isolates. Minimizing isolates and having an even color distribution is difficult, especially with only 16 rooms. With two colors, we were able to get good results with 31 corridors (maze six), but with three colors, it was not possible to find a maze over 2000 with only 31 corridors. Thus, we elected to use the largest number of corridors the grid structure underneath would allow, which was 42, given that we were not allowing multiple corridors between two rooms, even if the corridors were in opposite directions with different colors. The results demonstrate that it is possible to have a high scoring maze with three colors, 42 corridors and 16 rooms.

Due to having two players, small changes to the maze instance can result in large changes to the state graph, which makes local search more difficult for this type of maze. The instance scores are lower as well, and this is due to the color constraints and generally fewer notions of "doubling back" to increase the score.

It seems most natural for this type of maze to generate holes of substantial size; all of the results except for maze seven have substantial holes. Abbott noted in his description of this problem that it has an "abundance of loopiness" (Abbott, 1999, 37). Perhaps the two player color constraints, especially when the board is more monochromatic, enable this behavior.

The topologies of the top mazes are similar to the previous mazes - one large black and white hole combined with a whirlpool of the required vertices (see Fig. 5.8). Additionally,

an incorrect choice (one of two) from the starting vertex leads directly into the black hole, similar to Fig. 5.6.

5.5   Comparison to Abbott's Instances

5.5.1   Jumping Jim's Encore

The Jumping Jim's Encore instance Abbott provides has an overall score of 1438. It is an 8 x 8 instance, so it is the same size as the instance generated in the previous section. The shortest path of Abbott's instance is length 17, which is only 13.3% of the vertices; however, we do not apply the n^2 penalty for comparison purposes (this would result in a score of less than -14000). It has a path score of 1270, a branching score of 13, a decision score of 108, a reachability score of 144, and a dead end score of -105. It has 11 circled cells, five state changes, two doubly visited squares (one of them being the start square) and two u-turns. We can see, by this scoring metric, that the instances we have generated are generally higher-rated than Abbott's. This is due to several factors: only 58% of the vertices are reachable, 94% are reaching, and 94% are both reachable and reaching, so 5% of the vertices are isolates. The larger number of isolates combined with the small reachable portion of the maze limits the reachability score. There are no black or white holes although there are several long reverse false paths, indicated by the much larger portion of reaching vertices. The circling percentage is fine, and there are fewer swaps, doubling backs, and state changes when compared to our instances. Generally, the statistics are all lower, which is why the score of Abbott's instance is lower overall.

5.5.2   Apollo's Revenge

Abbott's Apollo's Revenge instance is 7 x 7 instead of 8 x 8, which makes a direct comparison with the results generated above more difficult because the state graph has only 97 vertices instead of 127. However, we generated several 7 x 7 instances for comparison purposes, and we found a similar result to Jumping Jim's Encore.

Abbott's instance has substantial dead ends (even more than the instances with 127 vertices), a low decision score of 26, no holes, a negative reachability score due to 15%

isolates, and generally poorer instance terms, for an overall score of 869.

By comparison, we generated instances with scores of over 1500, with half as many dead ends, sizeable holes, and decision and reachability scores over 400 and 100, respectively.

5.5.3 Grandpa's Transit Map

The Grandpa's Transit Map instance in Mad Mazes has 142 vertices, which is more than the ones we generated previously, but has a lower score of 1708. Similar to the jumping mazes, the shortest path is just over 35% of the vertices (36%), so we modified the restriction placed previously for easier comparison. This instance has a branching score of 24, a 265 reachability score, a dead end score of -65, and a decision score of 28. These traits are similar with some being better than some of the instances we generated except for the decision score. Abbott's instance does not have any large holes, and the backward decision score is only 28, compared to the forward decision score of 116. There is a whirlpool that contributes to the decision score when moving forward, but this whirlpool is close to the start and is vulnerable to back-tracing, hence the forward decision score is much higher than the backward. Additionally, although Abbott's instance has a large shortest path, many of the decisions are forced, that is, there is only one outgoing edge and no decision must be made.

5.5.4 Spacewreck

The Spacewreck instance in *Mad Mazes* has 379 vertices; however, it scores very poorly due to a large quantity of isolates. 30% of the vertices are reachable, 22.43% are reaching, and 40% are both reachable and reaching. This means that 60%, of the vertices, or $0.6 * 379 = 227$ vertices are both unreachable and unreaching, i.e. 60% of the maze is wasted space. This results in a large negative score because 227^2 is subtracted from the score. The forwards decision score is 194 due to a black hole; Abbott directly references this black hole in the description of the maze where he gives a sequence of moves and concludes with "alas, if you continue with this example, you'll find that both space travelers will soon be trapped in endless loops" (Abbott, 1990, 32). The black hole consists

of 12 vertices, and there are six entrances to it along the shortest path. However, the backward decision score is only two. Only a tiny portion of the maze is reaching, making it extremely vulnerable to back-tracing. Considering backwards movement from the finish, there are 25 options for the first move, but all of them are immediate dead ends except for the correct move. After this move, one can easily back-trace to the start because all of the decisions are either correct or immediately dead end after one move.

Figure 5.6: Maze one of the connections mazes - exceptional holes and a whirlpool combined with a large number of required decisions.
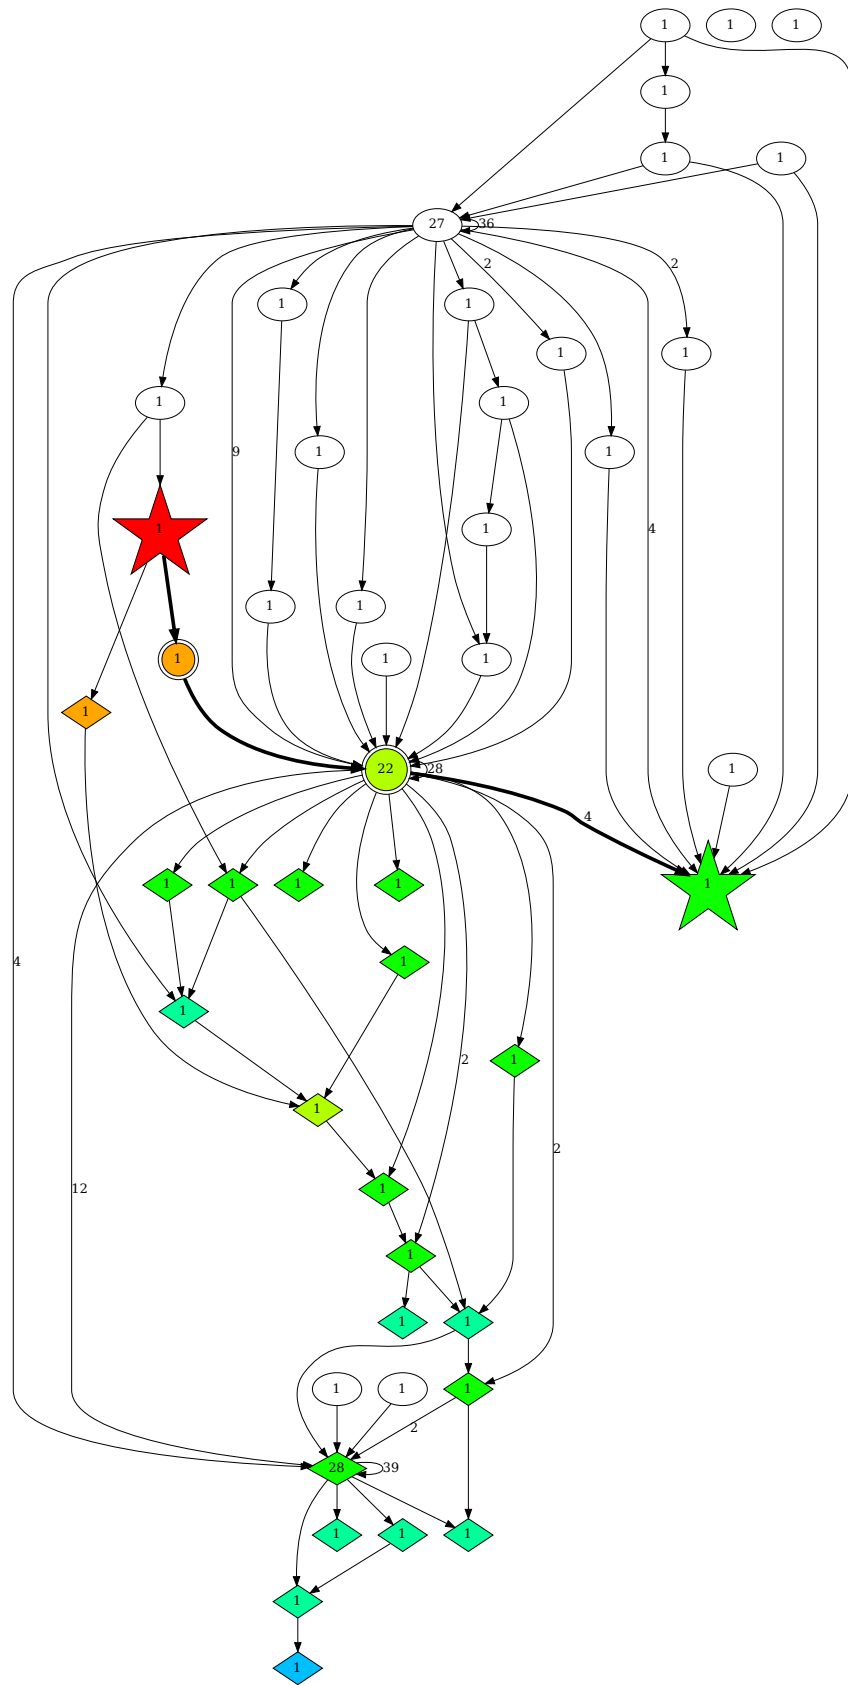
Figure 5.7: The best Spacewreck instance (three color).

Figure 5.8: The component graph associated with the best Spacewreck instance given in Fig. 5.7.

CHAPTER 6

CONCLUSIONS AND FUTURE WORK

6.1 Future Work

6.1.1 Human Maze Solving

We create mazes to be difficult for humans because they are simple to model and solve on a computer. Our approach consisted of creating objective functions and maximizing their value based on maze characteristics. However, these objective functions need to be validated. One possible approach is to simulate how a human would solve these mad mazes and compare the simulation results to the objective function.

There is a plethora of research on human methods to solve classic mazes, but none that we have found on mad mazes. Thus, we have looked into the classic maze methods to find potential similarities to mad maze solving methods. Determining human methods to solve mad mazes and extrapolating classical maze human interaction to logic mazes both need further research.

Zhao and Marquez (2013) identify two "stages" of human maze solving: exploration and guidance. In the exploration phase, the solver's eyes rapidly trace potential paths through the maze. Once a path has been chosen, the path is traced out (this study was performed digitally, so the guidance phase was the subject guiding the mouse along the chosen path). Complex mazes lead to often multiple stages of exploration and guidance (Zhao and Marquez, 2013, 6-8). On smaller mazes, some solvers attempted a brute force method, tracing paths through the maze continually until they found the solution, but this makes mistakes very costly and a more analytical approach was often used on larger mazes.

Karlsson used a DFS to simulate human maze solving, although he mentioned a DFS is not a perfect simulation because human choices at intersections are non-deterministic (Karlsson, 2018, 2). Although we are running this human-simulated algorithm on the state graph, it is as if we implicitly discover the state graph, as a human would not enumerate the entire state graph before solving the maze. We ultimately decided to use a randomized

DFS with many runs, and the average number of moves before reaching the solution was used as the difficulty metric. The results were somewhat similar to the objective functions that we created, but often were more ordered by the number of vertices in the state graph, which naturally leads to longer average DFS completion times. Additional work to identify a suitable algorithm that sufficiently models human state-space exploration could be used as another difficulty rating for state graphs.

It is possible for games such as Sudoku to be modeled as constraint satisfaction problems, and Pelánek (2011) has been able to model human Sudoku solving behavior with success. However, comparisons to games such as Sudoku are difficult in this context because they are fundamentally different problems. Sudoku is a constraint satisfaction problem and has been proven to be NP-Complete in the general case. An analysis of Sudoku difficulty presented by Pelánek showed an excellent correlation with human solving of 0.85 (Pelánek, 2011, 2-6), but additional work is necessary to determine if this could be mapped to logic mazes.

6.1.2 Additional Logic Mazes

This work could be continued for several other types of Mad Mazes presented by Abbott, such as Alice/Step Change Mazes, Rolling Cube Mazes, Theseus and the Minotaur, and more. Theseus and the Minotaur has been turned into a game and monetized as a mobile app (Abbott, 2010, 1). One could theoretically create such an app to test the difficulty of multiple different logic mazes with human players.

6.1.3 State Space Puzzles

Any game that can be modeled as a state space can be represented using this model, and its characteristics and traits rated by the scoring methods we have developed. This has potential for applications in the gaming industry (a multi-billion dollar industry that continues to grow). Modeling a game as a state space and identifying qualities the space must contain to be entertaining and challenging for humans to solve is a potential application. Generally, the state space should look like a state space of a maze, especially

in puzzle games. Many other games such as Sokoban can be represented as state spaces (Jarusek and Pelánek, 2010, 2) and rated in such a fashion as the logic mazes presented in this work.

6.1.4   Exploring Additional Optimization Techniques

There is an abundance of literature in the AI based searches; we used a relatively simple search method as a proof-of-concept. Other methods may produce better results with certain state graphs or mazes. Additionally, it may be possible to reduce the maze-generation problem to SAT or ILP, and then use an SMT solver such as Z3 or an ILP solver. This is another potential avenue for future research.

6.1.5   State Graph to Maze Instance

As noted in Chapter Four, it is theoretically possible to build a perfect state graph that has all of the necessary properties of the puzzle instance, and invert it back to the puzzle. Whether it is possible to formally characterize such a state graph in graph theoretical terms is itself an interesting research question. If the answer is yes, the next research task would involve developing algorithms to generate such state graphs.

6.1.6   Applications in Pedagogy: Puzzle-based Learning

Puzzle-based learning is a pedagogical strategy which teaches problem-framing and problem-solving for unstructured problems (Meyer et al., 2014, vi). The authors of this work claim these skills are generally lacking in today's students. A possible area of future work would involve a more formal pedagogical exploration of mad mazes as part of a puzzle-based learning strategy that assumes students have background in basic graph algorithms and concepts.

6.2   Conclusions

We have answered the research question of how to intelligently generate mad mazes instances: using state graph representations, objective functions, and local search. It is vital to consider not only the state graph's characteristics, but also the qualities of the maze instance itself to avoid undesirable properties such as monochromatic puzzle

instances that are designed to be poly-chromatic. Both stochastic local search and simulated annealing provided desirable results in the mazes that were generated. The most notable contribution of this work is that the same objective function can be used to rate any mad maze state graph, and potentially even outside mad mazes, to any game that can be modeled as a state graph.

REFERENCES

Abbott, R. (1990). *Mad Mazes: Intriguing Mind Twisters for Puzzle Buffs, Game Nuts and Other Smart People.* Adams Media Corporation.

Abbott, R. (1997). *Supermazes: Mind Twisters for Puzzle Buffs, Game Nuts and Other Smart People.* Prima Publishing.

Abbott, R. (1999, Dec). Logic mazes. https://www.logicmazes.com/dec99.html.

Abbott, R. (2010, Jul). Theseus and that pesky minotaur. https://www.logicmazes.com/theseus.html.

Buchin, K., M. Buchin, E. Demaine, M. Demaine, D. El-Khechen, S. Fekete, C. Knauer, A. Schulz, and P. Taslakian (2007, 01). On rolling cube puzzles. pp. 141–144.

Cooper, K., T. Harvey, and K. Kennedy (2006, 01). A simple, fast dominance algorithm. *Rice University, CS Technical Report 06-33870*.

Fisher, A. and G. Gerster (2000). *The Art of the Maze.* Seven Dials.

Hagberg, A. A., D. A. Schult, and P. J. Swart (2008). Exploring network structure, dynamics, and function using networkx. In G. Varoquaux, T. Vaught, and J. Millman (Eds.), *Proceedings of the 7th Python in Science Conference*, Pasadena, CA USA, pp. 11 – 15.

Jarusek, P. and R. Pelánek (2010). Human problem solving : Sokoban case study.

Karlsson, A. (2018). Evaluation of the complexity of procedurally generated maze algorithms.

Meyer, E. F., N. Falkner, R. Sooriamurthi, and Z. Michalewicz (2014). *Guide to Teaching Puzzle-based Learning.* Springer London.

Neller, T. W., A. Fisher, M. T. Choga, S. M. Lalvani, and K. D. McCarty (2011). Rook jumping maze design considerations. In H. J. van den Herik, H. Iida, and A. Plaat (Eds.), *Computers and Games*, Berlin, Heidelberg, pp. 188–198. Springer Berlin Heidelberg.

Pelánek, R. (2011, 01). Difficulty rating of sudoku puzzles by a computational model.

Russell, S. J. and P. Norvig (2022). *Artificial Intelligence: A Modern Approach* (4 ed.). Pearson.

Skiena, S. S. (2008). *The Algorithm Design Manual* (2nd ed.). Springer Publishing Company, Incorporated.

Zhao, M. and A. G. Marquez (2013). Understanding humans' strategies in maze solving. *CoRR abs/1307.5713*.